

March 19, 2023

Users Manual for the UEDGE Edge-Plasma Transport Code

T.D. Rognlien, I. Joseph, W.H. Meyer,
M.E. Rensink, and M.V. Umansky
Lawrence Livermore National Laboratory
Livermore, CA 94551

LLNL Report No.: LLNL-SM-846481

Operational details are given for the two-dimensional UEDGE edge-plasma transport code. The model applies to plasma and neutral fluids in a strong magnetic field. Equations are solved for the plasma density, velocity along the magnetic field, electron temperature, ion temperature, and electrostatic potential. In addition, fluid models of neutrals species are included or the option to couple to a Monte Carlo code description of the neutrals. Multi-species ion mixtures can be simulated. The physical equations are discretized by a finite-difference procedure, and the resulting system of algebraic equations are solved by fully-implicit techniques. The code can be used to follow time-dependent solutions or to find steady-state solutions by direct iteration. The code is written in Fortran with a very small amount of amount C. Calculations can be performed under Python or Basis code steering frameworks, which provide very similar functionality, with Python become the preferred method for new users.

Contents

1	INTRODUCTION	6
2	Source Code, Variable Descriptor Files, and Building an Executable	7
2.1	Source code	7
2.2	Variables	8
2.3	Compiling and loading a new executable	10
2.3.1	Building UEDGE for BASIS with MIO	11
2.3.2	Building UEDGE under BASIS with MMM	12
2.3.3	Building PYUEDGE for use with PYTHON	12
3	Basic Mechanics of Code Execution	17
3.1	Reading input and execution	17
3.2	Sample problems	19
3.3	Displaying results	19
3.4	Instantaneous output and run termination	20
3.5	Restarting from present solution	21
3.6	Creating and restarting from a BASIS portable PFB save file	21
3.7	Creating and restarting from a PYTHON-based portable PDB save file	22
3.8	Interpolating the solution to a new mesh and restarting	23
4	Grid Generation	25
4.1	Mesh generation using MHD equilibria	25
4.2	Non-orthogonal grids	29

	3
4.3 Adaptive-mesh capability	33
4.4 Adding poloidal cells near the x-point	36
4.5 Top-of-mesh/limiter option	37
4.6 Cartesian and cylindrical configurations	37
5 Running UEDGE in the Time-Dependent Mode	38
5.1 Setting simulation time and diagnostic output	39
5.2 Calculation of Jacobian	39
5.3 Accuracy	40
5.4 Boundary conditions for the time-dependent mode	40
5.5 Using the NKSOL solver in a time-dependent mode	41
5.5.1 Interactive mode	41
5.5.2 Controlling NKSOL time-dependence with scripts	41
6 Running UEDGE with a direct Newton iteration to steady state	43
6.1 Switches and diagnostic output	43
6.2 Preconditioning for the svrpkg="newton" case	44
6.3 Determining convergence trends and abort command	44
7 Running UEDGE with Krylov-Newton Iterations to Steady State	44
7.1 Switches and diagnostic output	44
7.2 Preconditioning options	45
7.3 Row and column scaling and rescaling	46
7.4 Pseudo-transient timestep	47

8	Boundary Condition Options	48
8.1	Specifying gas input and pumping on the side-walls	48
8.2	Other side-wall boundary condition options	50
8.3	End-plate boundary condition options	52
8.4	Boundary conditions at the core interface	54
8.5	Core boundary conditions with cross-field drifts	55
8.6	Sputtering boundary conditions for the gas species	55
8.6.1	physical sputtering by ions on plates	56
8.6.2	chemical sputtering by neutrals on side walls	56
8.6.3	physical and chemical sputtering by ions on side walls	57
9	Sources and Sinks	58
10	Flux-Limiting Transport Coefficients	59
10.1	Basic flux-limit expressions	59
10.2	Extensions to flux-limit models	60
11	Models for Neutral Gas	62
11.1	Inertial fluid and diffusive models for atoms	62
11.2	Options for temperature of neutrals	63
11.3	Inclusion of fluid molecules via the diffusive approximation	64
11.4	Coupling to Monte Carlo neutral codes	65
12	Models for Hydrogen Ionization, Radiation, and Recombination	67
12.1	Basic rate data tables available	67

12.2	Escape-factor model for Lyman- α radiation trapping	68
12.3	Calculation of various atomic rates available at the parser	69
13	Model for Impurity Radiation and Transport	70
13.1	Fixed-fraction model	70
13.2	Multi-species models	71
14	Specifying Anomalous Radial Transport Coefficients	72
15	Converting solutions from Full-Space to Half-Space & Vice Versa	75
15.1	Converting from full-space to half-space	75
15.2	Converting from a half-space to a full-space	77
A	Equations used for the UEDGE code	79
B	UEDGE topology conventions	85
C	A Primer on Input Files and Obtaining Initial UEDGE Solutions	87

1. INTRODUCTION

This report gives the operational details of how to use the **UEDGE** code. A brief description of the equations solved is included in the Appendix; for more details, the reader can refer to the reference list at the end of this report which include more information on the models and examples of the results obtained with the code. The most up-to-date version of this manual is available on the web site <http://www.mfescience.org> and then opening the links to Theory Group / Uedge Doc.

UEDGE is a two-dimensional (2D) fluid transport code for collisional edge plasmas. Its primary use has been for tokamak edge plasmas in Magnetic Fusion Energy devices, mainly tokamaks, although linear devices and spheromaks have also been modeled. **UEDGE** typically generates a curvilinear mesh based on the poloidal flux surfaces from an MHD equilibrium code such as **EFIT** or **TEQ**, but there are options for Cartesian meshes and cylindrical meshes as well.

The basic physics equations are taken from Braginskii [1], with the addition of *ad hoc* anomalous or turbulence-enhanced transport coefficients for the direction across the magnetic field; transport along the magnetic field is taken as classical with flux limits. A discussion of the rationale for this procedure is given in Ref. [2]. Also, an arbitrary number of ion species can be included (limited by computer memory and speed) which goes beyond Braginskii's model. Line-radiation loss from excitation, ionization, and recombination is incorporated into the electron energy equation. Neutral gas is described by fluid models or by coupling to a Monte Carlo code.

At its inception in 1992, **UEDGE** used a set of basic physics equations and finite-differencing similar to that in the original **B2** transport code [3,4]. However, **UEDGE** uses a fully implicit procedure known as a modified Newton iteration to solve all of the equations simultaneously rather than the semi-implicit **SIMPLE** algorithm used in **B2**. Overviews of the fully-implicit method used in **UEDGE** are given in Refs. [5,6]. In addition, **UEDGE**, now includes a detailed fluid neutral model with a parallel momentum equation, calculation of the electrostatic potential with $\mathbf{E} \times \mathbf{B}$ and ∇B drift effects, and a nonorthogonal mesh to conform to shaped

divertor surfaces.

References 7–25 give many details of the models used in UEDGE and some applications for edge-plasmas in fusion devices. The list is not intended to be exhaustive, but to serve as beginning bibliography for those seeking more detail.

2. Source Code, Variable Descriptor Files, and Building an Executable

2.1. Source code

The source code is maintained in a CVS archive on the MFE Program network at LLNL. Precompilation processing can be done on any workstation that can access to this network, and others can obtain the necessary files by contacting the authors. Compilation and loading is done on the machine where UEDGE is to be run. Details of this procedure are given below in Sec. 2.3.

UEDGE is divided into ten BASIS packages with different functions. For example, the finite-differenced physics equations are contained in the package **bbb**, the grid generation routines are in packages **flx** and **grd**, and the (**pfb**) package (provided by BASIS if used) allows reading and writing data in a portable format PFB save file. An alphabetical list and summary of all of the packages (each having its own directory) and several important auxiliary directories is as follows:

aph	# calculates atomic cross-sections for hydrogen
api	# calculates atomic cross-sections for impurities
bbb	# sets up the finite-difference physics equations and routines
	# needed by the linear algebra solvers - the heart of UEDGE
com	# contains routines and variables needed by various packages
dce	# distributed computing package - not generally needed
doc	# documentation including UEDGE manual, uedge.man

```

dst      # DUSTT module track/ablate dust; also idf,psi (UCSD)
idf      # import data from file; used only by dst (DUSTT) (UCSD)
psi      # plasma/surface interact; used only by dst (DUSTT (UCSD)
flx      # calculates the magnetic flux surfaces for mesh
grd      # calculates second mesh coordinate and constructs the mesh
in       # various diagnostic routines (not a package)
ncl      # NCLASS modular for neoclassical transport (ORNL)
scripts  # BASIS scripts for finding files (not a package)
svr      # linear algebra and temporal integration routines
test     # a few test cases (not a package)
wdf      # calculates data needed by the DEGAS2 neutral M.C. code

```

2.2. Variables

All of the variables within UEDGE, together with a one-line description of their meaning, are listed in files called the variable descriptor files which are used in conjunction with either the BASIS or PYTHON scripting systems that are used to drive/steer UEDGE. Generally, there is one such file for each package listed above, with the name of the file being `package_name.v` in the directory `package_name`. For example, variables for the physics equations and routines for the numerical Jacobian in UEDGE are included in the variable descriptor file called `bbb.v` in directory `bbb`. In both BASIS and PYTHON versions, all variables include the prefix from the package in which they are defined, *i.e.*, most plasma physics variables have the `bbb.` prefix. However, an important difference between running under PYTHON compared to BASIS is that PYTHON requires you to use the prefix, whereas BASIS does not (although it allows prefixes to be used; if multiple names exist and none is specified, it uses the name associated with the package highest on the “stack”, *i.e.*, most recently used).

The variables are combined into groups to make the process of identification easier. Also, if you know the name of the variable while running UEDGE, at the UEDGE> prompt, you

may type `list xyz`, and you will receive the information about that variable included in the variable descriptor file. Also, if you know the name of the group, you can type `list groupname`, and receive a listing of that group from the variable descriptor file.

The primary plasma fluid variables used in the code are as follows (where the suppressed prefix of `bbb.` must be used in the PYTHON version, *e.g.*, `bbb.ni`):

<code>ni(0:nx+1,0:ny+1,nfld)</code>	Ion dens. [m^{*-3}], <code>nfld=species index (default=1)</code>
<code>up(0:nx+1,0:ny+1,nfld)</code>	Parallel ion flow velocity [m/s]
<code>te(0:nx+1,0:ny+1)</code>	Electron temperature [Joules]
<code>ti(0:nx+1,0:ny+1)</code>	Ion temperature [Joules]
<code>ng(0:nx+1,0:ny+1,ngsp)</code>	Neutral gas density [m^{*-3}], <code>ngsp=species index</code>
<code>phi(0:nx+1,0:ny+1)</code>	Electrostatic potential [Volts]

Note that SI units (meters, volts, sec, *etc.*) are used throughout UEDGE, and scaling (normalization) is applied to the final ODE's and at the linear algebra stage (row and column scaling). The first two indices for each variable correspond to mesh indices (`ix, iy`), where `ix` is the poloidal index beginning at the inner divertor plate for a tokamak, and `iy` is the radial-like index beginning at the core boundary or the private-flux-region wall.

The primary plasma fluid variables are used to evaluate the spatial derivatives and source terms in the PDE's. However, the variables that are passed to the ODE and Newton solvers are somewhat different and normalized. These are as follows:

For density:	$ni(i)/n0(i)$, where $n0(i)$ is an input constant
For velocity:	$ni(i)*up(i)/[n0(i)*cs]$, where cs is a constant ion-acoustic speed
For Te:	$ne*te/(n0(1)*tnorm)$, where $tnorm$ is a constant, $ne = \text{elec. dens.}$
For Ti:	$ne*ti/(n0(1)*tnorm)$, where $tnorm$ is a constant, $ne = \text{elec. dens.}$
For ng:	$ng(i)/n0g(i)$
For phi:	$ev*phi/tnorm$, where $ev=1.6e-19$ is the electron charge

The conversion from plasma variables to ODE variables is done in subroutine `convrs`; the reverse conversion is done by subroutine `convert`.

The ODE variables are stored in a 1-D vector call `yl`, starting at `ix=0`, `iy=0`. The variables at that point are stored in the order listed above, then the poloidal index, `ix`, is incremented until the `ix=nx+1` boundary is reach, then `iy` is incremented by unity, and the process repeated. There are index arrays that allow the user to determine the index `ieq` of `yl(ieq)` that corresponds to a variable at a given `(ix,iy)` location on the grid; these arrays are defined as follows:

<code>idxn(ix,iy,ifld):</code>	<code>yl</code> <code>ieq</code> index for <code>ni(i)/n0(i)</code>
<code>idxu(ix,iy,ifld):</code>	<code>yl</code> <code>ieq</code> index for <code>ni(i)*up(i)/[n0(i)*cs]</code>
<code>idxte(ix,iy):</code>	<code>yl</code> <code>ieq</code> index for <code>ne*te/[n0(1)*tnorm]</code>
<code>idxti(ix,iy):</code>	<code>yl</code> <code>ieq</code> index for <code>ne*ti/[n0(1)*tnorm]</code>
<code>idxg(ix,iy,igsp):</code>	<code>yl</code> <code>ieq</code> index for <code>ng(i)/n0g(i)</code>
<code>idxphi(ix,iy):</code>	<code>yl</code> <code>ieq</code> index for <code>ev*phi/tnorm</code>

There are also two arrays that the give the `ix` and `iy` indices for a given `yl` index `ieq`:

<code>igyl(ieq,1):</code>	<code>ix</code> poloidal index for ODE variable <code>ieq</code>
<code>igyl(ieq,2):</code>	<code>iy</code> radial index for ODE variable <code>ieq</code>

2.3. *Compiling and loading a new executable*

The procedure for compiling and building the **UEDGE** code is outlined below:

Begin in your home directory, and un-tar the files containing **UEDGE** by the command

```
tar xvf uedge_Vx.x.tar
```

where `Vx.x` gives the **UEDGE** version number. See the Appendix to correlate version numbers with dates.

Set the following environmental variables if you are using **BASIS** to help build and run **UEDGE**: note that `BASIS_ROOT` and `NCAR_G_ROOT` will depend on where the system ad-

administrator has stored BASIS, NCAR graphics, and PACT (see <http://pact.llnl.gov>) on your computer system

```
setenv UEDGE_SCRIPTS ~/uedge/scripts
setenv BASIS_ROOT /usr/local/nbasis
setenv NCARG_ROOT /usr/local
setenv PACT /usr/local/pact
setenv OPT '-native -O3'
set path = ($BASIS_ROOT/bin $path)
```

There are now two options for building UEDGE with BASIS, with the preferred new method utilizing the MIO system; the older method uses MMM as described in the next two subsections.

If instead you are building the PYTHON version of UEDGE, most of the required paths are set in the configure file `uedge/Python/config/ALL.config` as described in more detail in Sec. 2.3.3. Pyuedge will run without using PACT for creating and reading PDB format files for restart, but this is a very convenient feature. Thus, to use this option described below, you should install PACT (see <http://pact.llnl.gov>), and continue to Sec. 2.3.3.

2.3.1. *Building UEDGE for BASIS with MIO*

MIO is now distributed with the BASIS source code and needs to be installed on your machine. You can then learn some more about MIO by typing `man mio`. To build UEDGE, go to the `~/uedge/builder` directory and issue the following commands:

```
dsys config [-o] machine-dependent-file
dsys build
dsys load
dsys test
```

Here `[-o]` is the option to optimize (or `-g` for debug) and `machine-dependent-file` is one of the file names in the `builder/std` directory that contains MIO switches appropriate for different

computer architectures, compilers, and options; for example, on linux machines, this file name often begins with (or is) `linux`. Thus, this is the file that the user may need to edit for their specific machines. The `dsys test` command shown above runs a set of simple tests with the new executable to confirm that the executable produces the same result for some base cases. If there is a physics change or modification of the numerical scheme, the reference pdb file needs to be updated in `test/level.2/ref` directory; if this pdb file is removed, the test script will properly add a new one after obtaining the present solution.

The executable is located in `~/uedge/dev/machine-dependent-name/bin/xuedge`. Here `machine-dependent-name=lnx-2.1-i32`, for example, on some LINUX machines. The xuedge executable produced by MIO should be functionally equivalent to the executable produced the old way with MMM, but it is stored in a different location.

2.3.2. *Building UEDGE under BASIS with MMM*

Go to the `~/uedge` directory and issue the following to generate the required make files

```
mmm -ezn -rl
```

where `ezn` is the graphics package and `rl` is the readline library that allows line editing and line recall. This command generates the makefiles needed to compile and load UEDGE.

Finally, type the following line to compile and load the executable code `xuedge`

```
gmake all xuedge
```

The executable `xuedge` will appear in your `uedge/LINUX` directory if you are on a LINUX system and more generally, in `uedge/$CPU`.

2.3.3. *Building PYUEDGE for use with PYTHON*

It is possible to build UEDGE without BASIS, but still having a scripting shell driven by PYTHON. Contributing to the development of PYUEDGE have sequentially been Dave

Grote, Brian Yang, Tom Rognlien, and Lynda LoDestro. The only external source that you may need is the preprocessor MPPL, which converts the .m files of the UEDGE source code to FORTRAN .f files. The MPPL source and instructions can be obtained on the web from

<http://w3.pppl.gov/rib/repositories/NTCC/files/mppl.tar.gz>

The UEDGE source provides the Mac_scripts and PyMAC scripts for convert the variable descriptor files (.v) to FORTRAN. The content of this section is included in updated form in the UEDGE source file uedge/README_Pyuedge, and the user is encouraged to read this for any update differences to the following discussion.

To construct the PYTHON version of UEDGE, follow these steps:

Checking out and building the needed modules:

1. If you have access to the LLNL/FEP cvs archives, then


```
setenv CVS_RSH=ssh
```

```
setenv CVSROOT :ext:<login-name> hrothgar.llnl.gov:/usr/local/cvsroot
```

 - (a) If BASIS is installed on the target computer, then


```
cvs co [-P] pyUtils uedge # -P empty dirs not checked out
```
 - (b) or, if BASIS is not installed, add Mac_scripts to cvs co, i.e.,


```
cvs co [-P] pyUtils uedge Mac_scripts # -P empty dirs not checked out
```
2. or, if you get your copy of UEDGE from a tar file
 - Set up a UEDGE directory and untar the file that should include the directories pyUtils, uedge, and Mac_scripts(if BASIS is not on system)
3. Then, for either (1) or (2):
 - (a) If BASIS is installed on your system (likely at LLNL, PPPL) then
 - Check that compiler paths, etc. are valid on your machine in the file uedge/Pyuedge/config/*.config : Defaults are in config/ALL.config.

- Users must explicitly set the path variable `TOPALL` in the file `uedge/Python/Makefile` to the parent directory of `uedge` and `pyUtils` directories
- Then


```
cd uedge/Pyuedge
gmake make.all # builds pyUtils and Pyuedge
```
- A shared object file is produced as `uedge/Pyuedge/yoursys/uedge.so`, where `yoursys` is the architecture variable `ARCH`, e.g., `Linux`
- To load `uedge.so` into a python session, startup python using the same version as defined in


```
uedge/Pyuedge/config/${ARCH}.config
```

 and read in the `uedge` startup script


```
uedge/Pyuedge/uedge_startup.py, i.e. python
>>>execfile('uedge_startup.py')
```

(b) If `BASIS` is not installed on your system

- `MPPL` must be installed first. `MPPL` can be downloaded from <http://w3.pppl.gov/rib/repositories/NTCC/catalog>
- Check the directory structure from (1) or (2) above; must include 3 directories: `pyUtils`, `uedge`, and `Mac_scripts`
- Check that compiler paths, etc. are valid on your machine in the file `uedge/Pyuedge/config/*.config` : Defaults are in `config/ALL.config`.
- Users must explicitly set the path variable `TOPALL` in the file `uedge/Python/Makefile` to the parent directory of `uedge` and `pyUtils` directories
- Then


```
cd uedge/Pyuedge
gmake make.all # builds pyUtils and Pyuedge
```
- A shared object file is produced as `uedge/Pyuedge/yoursys/uedge.so`, where

yoursys is the architecture variable ARCH, e.g., Linux

- To load uedge.so into a python session, startup python using the same version as defined in

```
uedge/Pyuedge/config/${ARCH}.config
```

and read in the uedge startup script

```
uedge/Pyuedge/uedge_startup.py, i.e.
```

```
python
```

```
>>>execfile('uedge_startup.py')
```

4. Additional useful commands/information

Other gmake commands:

- gmake # builds only Pyuedge
- gmake clean # cleans only Pyuedge
- gmake clean.all # cleans both pyUtils and Pyuedge

Key make-variables for customizing directory locations:

CONFIG - path to the config directory. It is set in Pyuedge/Makefile

PYUTILS - path to the pyUtils directory. It is set Pyuedge/Makefile (overrides the default, which is set in config/ALL.config)

HasBASIS - determines which mac and mppl are used; defaulted in config/ALL.config.

If HasBASIS is defined, the following must be set:

BASIS_ROOT - path to basis; mac and mppl from basis are used.

If HasBASIS is undefined, the following are needed:

MAC - The mac utility executable; defaulted in config/ALL.config.

MPPL_ROOT - path to MPPL. There is no platform-independent default set up (gmake will crash if MPPL_ROOT is not found).

PYTHON_ROOT, PYVERS - path and python version

Compiler, etc., paths and options are defaulted in config/ALL.config. The platform-dependent settings are in the appropriate config/*.config file.

To run a simple test case (which automatically reads in `uedge/Pyuedge/uedge_startup.py`), do the following in `uedge/Pyuedge`

```
python
>>>execfile('rdtest_linux.py')
```

The session should look something like the following:

```
hrehtric:rognlien(LINUX)> pyuedge
Python 2.2 (#3, Mar 11 2002, 13:29:17)
[GCC 2.96 20000731 (Red Hat Linux 7.1 2.96-85)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> execfile('rdtest_linux.py')
```

You type the last line to read the file `rdtest_linux.py`, which contains input parameters and the call to the driver routine `bbb.exmain`. This case solves for four plasma variables (density - `bbb.ni`, parallel velocity - `bbb.up`, electron temperature - `bbb.te`, and ion temperature - `bbb.ti`), which are advanced time-dependently, nearly to steady state. The four variables are printed at the end of the run. Also note that to execute **UEDGE** under **PYTHON**, that you must call the subroutine **exmain** as `bbb.exmain()`.

Note that in **PYTHON**, all variables and subroutines carry the name of the package with which they are associated, e.g., the ion density is `bbb.ni(ix,iy)` rather than just `ni(ix,iy)` in the **BASIS** version. In fact, **BASIS** also names `ni` \rightarrow `bbb.ni`, but the `bbb.` is suppressed unless the user explicitly includes it in commands.

Also note that `rdtest_linux.py` has an initial command `execfile('uedge_startup.py')` that reads the file `uedge_startup.py`, which imports the required packages into the `pyuedge` session.

Again be sure that the python version you run here is the same version as was used in the build (`(PYTHON_ROOT)/bin/python(PYVERS)`), where `PYTHON_ROOT` and `PYVERS` are as defined in your `uedge/Pyuedge/config/*.config` file.

To modify `uedge` sources and recompile, note that there is no `uedge` source-code under

uedge/Pyuedge except for the file Pyuedge/Fsor/ex_das_isa.f. Make changes in the uedge source-code directories parallel to Pyuedge, *e.g.* uedge/bbb or uedge/com, *etc.*; then return to Pyuedge and type gmake.

3. Basic Mechanics of Code Execution

3.1. Reading input and execution

(Note: this section was written to describe running UEDGE under BASIS although the basic steps have close counterparts when using the PYTHON version.)

The executable for UEDGE is typically called `xuedge`, or sometimes just `uedge` on NERSC machines. The build and loading procedure is described in Sec. 2.3, which results in the executable being found in `~developer/uedge/SOL`, where `developer` is the name of the person who created the executable and `SOL` refers to the SUN Solaris system. On other workstations, `SOL` will be replaced by something like `HP700` for HPs or `AXP` for DEC's, etc.

Before running the BASIS version of UEDGE, or `xuedge`, you need to set two environmental variables if you have not already done so for the compilation and loading of UEDGE described in Sec. 2.3. If you are running on the LLNL MFE SUN solaris system, you can just set two environmental variables: `UEDGE_SCRIPTS` to `/home/rognlien/uedge/scripts` and `UEDGE` to `/home/rognlien/uedge`. Then you need not worry about the `aphdir` and `uedge_path` files. If you are on another system, or want to read your own atomic physics and script files, place the contents of `uedge/in/aph` and `uedge/in/api` in some directory, and then set the environmental variable `UEDGE` to point to them. Likewise, copy and then change what is in `uedge/scripts` and set up the environmental variable `UEDGE_SCRIPTS` to point to that directory.

The `aphdir` file tell UEDGE where to locate the hydrogen atomic rate files, and on the LLNL MFE SUN solaris system, it should contain the line

```
aphdir = "/mfe/theory/Uedge/Ver_XX/uedge/in/aph"
```

where the `XX` is the most current version number shown in the Uedge directory. The second

file, `uedge_path`, tells UEDGE where to find various diagnostic files. Again, on the LLNL MFE SUN solaris system, it should contain the line

```
call pathadd("/mfe/theory/Uedge/Ver_XX/uedge/in")
```

One can add other lines to tell UEDGE to search other paths (directories) in response to a command issued from the parser.

If you are running the PYTHON version of UEDGE, you need to set the character variable `aph.aphdir` in your input file to the path of the directory containing the hydrogen atomic data. For example, this may be

```
aph.aphdir = "/mfe/theory/Uedge/Ver_XX/uedge/in/aph")
```

To run a case, simply type `xuedge` (or `uedge` if you are using the public version on the NERSC system). The executable will then automatically read a default input file called `.basis`, if it exists in the directory (generally not used, however). After UEDGE has read the `.basis` file, it comes back with the prompt `UEDGE>`, at which point you need to type a command. Typically, you will read a second file which contains the input settings for various control variables and switches; we usually use the convention that such input files begin with the letters `rd` (which stands for “read”), but this is not necessary. Thus, you will type something like `read rddata2`. You may also start the executable and read the input file on a single line with the command `xuedge read rddata2`. Some example files can be found in the location noted below in Sample Problems section. You can also change any variable at the prompt by typing, say, `runtim=1.e-10`, to modify `runtim`.

Once all the variables have been set, you execute the code by typing `exmain`. What this does is tell the BASIS system to execute the subroutine named `exmain`, which calls all the appropriate subroutines to execute a full run. It is also possible under BASIS to call other subroutines independently, but this is usually only done for debugging purposes. Thus, a typical session with UEDGE will look as follows:

```
xuedge
```

```

UEDGE> read rddata2
UEDGE> exmain
UEDGE>

```

The last prompt means the code has successfully completed the run and is ready for more input. If you want to stop, just type **end**. You can also display the results without exiting from the BASIS session.

3.2. *Sample problems*

There are a set of sample problems that can be run. For example, on the LLNL MFE LINUX system, one can go to the directory `~rognlien/uedge/Applic/Examples`. It is best to check with one of the current users to learn more details of these examples. Note that these examples pertain to running UEDGE under BASIS, but there is a very close correspondence to running under PYTHON.

3.3. *Displaying results*

You may want to list any variable or array to see what the solution looks like. For example, the 2-D electron temperature is stored in the array `te(0:nx+1,0:ny+1)` in units of Joules. You may specify the number of places to be displayed after the decimal point by typing `fuzz=2`, for example. Thus, the following sequence will print out the electron temperature array in electron volts:

```

fuzz=2
te/ev

```

where `ev=1.6e-19 Joules/eV` is a variable in the code for converting from Joules to electron volts. If you want to list some other variable to 4 decimal places you must first type `fuzz=4`, which will hold until another `fuzz=` statement is typed.

You can also do plots of arrays to look at the present solution. To plot the function $y(x)$, type `plot y,x`. To do a contour plot of $z(x,y)$, type `plotz z,x,y`.

All of the commands useable at the `UEDGE>` prompt are in the **BASIS** command language, which is extensive and powerful yet easy to learn. The **BASIS** manual can be read and searched with a web browser (see the URL <http://xfiles.llnl.gov/basis/>). The plotting package **EZN** is also part of **BASIS**.

3.4. *Instantaneous output and run termination*

While running a time-dependent problem with any solver (except `lsode`, which is not a currently supported option), you may obtain information on the present status of the solution by typing `s` or `status` and a return. In order to do this, you must first set the switch `iskaboom=1`; this option does not work on some platforms - notably, PCs running LINUX and possible others. If the code appears to “hang” without doing anything, one cause can be `iskaboom=1` on a platform that will not support this options; try setting `iskaboom=0`. Also, if you run `xuedge` in the background, you need to set `iskaboom=0`.

If `iskaboom=1` and you type `s`, you will receive lines of output giving the total number of function and preconditioner evaluations, `nfe` and `npe`, respectively, and the `yl` index `ieq` (`imxtstep`) of the equation that is restricting the time step. A second index, `imxnewt`, gives the `ieq` of the equation that is requiring the Jacobian to be reevaluated, if that is limiting the time step. In addition, it gives the total simulation time and the present time step, `dt`. A third line gives detail data on the error estimates from the ODE solver: `bigts` is for the time integration and `dsm` is for the Jacobian.

If you want to abort the present simulation and return to the `UEDGE>` prompt, type `ctrl-c` and a return. At this point, **BASIS** will give the prompt `DEBUG>`; you now may query `UEDGE` for variable values, etc., and then type `cont` if you want to continue the calculation. If you type `abort`, to the `DEBUG>` prompt, you will return to the `UEDGE>` prompt. You can do this for `DEBUG>` during either a time-dependent solution or a Newton iteration. If you then restart after changing some parameters, the code initializes itself as though the

aborted run had not taken place.

3.5. *Restarting from present solution*

If you are still in an active **BASIS** session, you can set `restart=1` to use the present solution as initial conditions for the next execution. You can also double the grid in both x and y directions by being sure that `newgeo=1`, `restart=1` and by doubling `nxleg`, `nxcore`, `nycore`, `nysol`, and `nxomit`. The present solution is then interpolated to the finer grid as the initial conditions for the new run. This is conveniently done by reading a file called `double` with **BASIS** (i.e., type `read double`) that has the necessary parameter adjustments in it. You can double the mesh separately in the radial or poloidal direction, or incrementally; see Sec. 3.8 for more details.

3.6. *Creating and restarting from a BASIS portable PFB save file*

BASIS has a very useful facility to save any variables you wish to a portable file that can be read in a subsequent session. To do this, type `create sfile`, where `sfile` is some name you choose for the data file. Then you can save any variables you want by typing `write x,y,te`. This command can be repeated; to stop the writing to this file, type `close`. The minimum data set you must save in a portable file to make a restart is the set of plasma variables; you should also save any special variable settings. An example for the minimum saving procedure is as follows:

```
create stest1
write nis,ups,tes,tis,ngs,phis
close
```

UEDGE uses the convention that the plasma variables used for restarting all have the letter “s” appended to them.

To use the portable file, you must be sure that you have the same grid size as that of the

saved data. You then execute **UEDGE** as for a normal run by reading input files (but don't type **exmain** yet). Then give the following commands:

```
allocate
restore stest1
restart=1
```

Here, **allocate** generates the appropriate arrays through dynamic allocation. At this point, one may type **exmain** to begin the run with the values saved previously in **stest1**. It is convenient to save the final state of converged runs in case you want to restart from them at some later date. Also, the portable files can be moved (using the binary mode of **ftp**) to computers using different numerical representations and used to continue runs.

Before version 10.0 of **BASIS**, a capability was present to create savefiles using commands **save to**, **save**, and **save off**; users of more recent versions do not need to concern themselves with this difference. The portable file capability (PFB or PDB) should be used now to create new files, but the user may want to use a version of **UEDGE** created with version 9.11 of **BASIS** to convert from the savefile format to the

portable format. Use the **readb** command to read a savefile just as you would use **restore** to read a portable file. Then use **create**, **write**, and **close** to save the data in a portable file.

3.7. *Creating and restarting from a PYTHON-based portable PDB save file*

If you are using **PYUEDGE** rather than the **BASIS** version, you can still create and restore saved solutions from PDB files. To do so, you need to have installed libraries associated with **PACT** and the scripts from **PyPDB**. Information in this section came from Dave Grote, LBL. **PACT** is a suite of tools developed at LLNL that allows creating and reading self-describing pdb data files (as used by **BASIS** above). **PACT** can be found at <http://pact.llnl.gov>. **PyPDB** are scripts that give **PYTHON** access to these tools.

After installing PACT and PyPDB, the user can create a PDB save file under a pyuedge session as follows:

```
ff = PW.PW("pf_case1")
ff.nis = bbb.nis
ff.ups = bbb.ups
ff.tes = bbb.tes
ff.ngs = bbb.ngs
ff.phis = bbb.phis
ff.close()
```

To read such a file back into PYEDGE, the user can type

```
restore('pf_case1')
```

(single or double quotes should work), or

```
ff = PR.PR("pf_case1")
bbb.nis = ff.nis
bbb.ups = ff.ups
bbb.tes = ff.te
bbb.tis = ff.ti
bbb.ngs = ff.ng
bbb.phis = ff.phi
ff.close()
```

3.8. *Interpolating the solution to a new mesh and restarting*

UEDGE has three different linear interpolating options; these are controlled by the switches `isnintp` and `isgindx`. If `isnintp=0`, `isgindx` is immaterial, and the “old” interpolator is used that only allows the users to double the mesh in each direction. That is, if you are starting from

a solution with a certain set of grid indices `nxleg(1,1)`, `nxleg(1,2)`, `nxcore(1,1)`, `nxcore(1,2)`, `nysol(1)`, and `nycore(1)`, you must double all of these. This interpolation is rather crude in that it assumes the mesh is uniform in each direction, but it works surprisingly well. However, doubling the mesh in each direction is sometimes too large a change for the Newton method to converge on the finer mesh. The next two methods allow an arbitrary change in the number of mesh points.

If `isnintp=1` the user may increase or decrease the mesh by any amount in either direction. For this case, two options are available: `isgindx=1` and 0. For `isgindx=1` (default and recommended setting), the interpolation occurs in index space as though the mesh is uniform. This case is thus similar to the `isnintp=0` option discussed above. However, there is a difference (other than the arbitrary mesh change) in that here values are not interpolated across the separatrix or across the radial cut through the x-point, but are rather extrapolated at these locations. The reason for doing this across the separatrix is that linear interpolation often does a poor job because of the abrupt change in variables there; it is done across the radial cut only for simplicity of the algorithm. Treating the region above and below the separatrix independently (and on either side of the cut) allows one to add extra mesh points in either region without perturbing the other.

For `isnintp=1` and `isgindx=0`, the code does a linear interpolation using the actual (normalized) mesh which is not uniform or rectangular. This scheme does not seem to outperform the simpler index-based algorithm, and sometimes has trouble finding the appropriate mesh points for performing the interpolation. The message `***** grdinty cannot find straddling grid...` is printed if the algorithm fails to find the appropriate mesh points. In that case, switch to `isgindx=1`.

It is possible to interpolate from a save-file solution, even if the code does not converge to this case first. One needs to generate the old mesh and then type `call gridseq` from the UEDGE prompt (the BASIS parser). Alternatively, use a very loose tolerance for `svrpkg="nksol"`, say `ftol=1.e10`, so the code will think it has converged after one iteration. The mesh may then be changed without any extra call to `gridseq`.

4. Grid Generation

4.1. Mesh generation using MHD equilibria

The grid generated in UEDGE uses routines that are an extension and modification of a grid generator developed by M. Petravic at PPPL [26]. Data on the location of poloidal magnetic flux surfaces are generated by one of various MHD equilibrium codes (*e.g.*, TEQ, EFIT), and then read into UEDGE via two files. These files must be named `aeqdsdsk` and `neqdsdsk`, and the switches `mhdgeo=1` and `gengrid=1` must be set. If `gengrid=0`, UEDGE reads in a file with the grid information already in it called `gridue`. The `gridue` file can be generated by a previous UEDGE run. For large problems, precomputing the `gridue` file and using `gengrid=0` can save considerable storage. One can generate the grid in UEDGE without running a complete problem; just type

```
call flxrun
call grdrun
```

and the file `gridue` will be generated. This call sequence is done automatically if you execute a full problem by typing `exmain`.

There are a number of configuration options available for grid generation with `mhdgeo=1`; we mention a few here. Set `geometry="snul"` for lower-single-null configurations; if there is a secondary (upper) x-point, the outermost flux surface will be adjusted by the code so as to exclude the second x-point. Set `geometry="uppersn"` for upper-single-null configurations; in this case the observer reference frame is simply transformed so the configuration appears as a lower-single-null. Set `geometry="dnbot"` to generate a half-mesh that contains a lower x-point and a symmetry boundary at the midplane; this option can be used with any `eqdsdsk` that contains a lower x-point (either lower single-null or double-null configurations). Set `geometry="dnull"` to generate an up/down symmetric double-null mesh that includes both x-points; the lower half-mesh (as described above for `geometry="dnbot"`) will be exactly duplicated in the upper half to form the complete mesh; this option can be used with any `eqdsdsk` that contains a lower x-point (either lower single-null or double-null configurations).

To generate a mesh that includes both x-points of a non-symmetric double-null configuration, one generates separate meshes for the upper and lower halves (as in `geometry="dnbot"`), and combines these to form a complete mesh for the configuration; see the example given below.

The construction of a full double-null mesh is illustrated in the script `makemesh_dRsep.m1` and associated files. These files are stored in the `uedge/test/rensink/dnull` subdirectory of the UEDGE source code distribution. This example uses a very small number of grid points to more clearly exhibit certain features of the mesh; for a realistic simulation one would use much finer resolution in both the radial and poloidal directions. The construction procedure consists of the following eight steps which are identified in the script that implements this example:

1. read the eqdsk to identify x-points and separatrices
2. specify the radial distribution of flux surfaces
3. construct the mesh for the upper half of the configuration
4. construct the mesh for the lower half of the configuration
5. combine the two halves of the mesh
6. define guard cells at the divertor target plates
7. compute magnetic field values for each cell
8. write out the `gridue` file

It is important to note that for any up/down asymmetric double-null configuration, the radial mesh option `kymesh=0` MUST be used. Also, steps 3 and 4 can be modified to create either orthogonal or non-orthogonal meshes, as desired.

It is also possible to simulate the outer half of a single null or the lower, outer quadrant of a double null where reflection boundary conditions are used at the left boundary and no flux is allowed through the outer (`ix=ixpt2`) cut at the x-point. To do this latter type of geometry, set these variables:

```

nxomit      # Number of poloidal grid points to omit before setting
             # reflection boundary condition. Files double, etc.
             # automatically change nxomit to correct value
isfixlb = 2  # Sets reflection bc at ix = nxomit and no-flux bc at
             # ix = ixpt2 (outer x-point cut)

```

The radial distribution of the mesh is controlled by the following input parameters. The poloidal flux, ψ , is normalized to be unity on the separatrix and the radial boundaries of the mesh are specified by:

```

psi0min1    :flux at the inner core boundary, 0.98 is typical
psi0min2    :flux at the private-flux region boundary, 0.98 is typical
psi0sep     :flux very near the separatrix, use 1.00005
psi0max     :flux at the outer wall boundary, 1.07 is typical

```

The number of radial cells in various regions is:

```

nycore(1)   :number of radial cells in core (and private-flux) region
nysol(1)    :number of radial cells in scrape-off layer

```

The radial distribution of mesh flux surfaces is controlled by the input parameter `kymesh`. For `kymesh=1` (the default value) analytic forms are used to radially distribute the normalized flux values between `psi0min` and `psi0sep`, and between `psi0sep` and `psi0max`; the input variable `alfcy` provides some additional control over the radial distribution. One can cause the flux surfaces to cluster near the separatrix by making the variable `alfcy` somewhat larger than unity (2-3 is typical); if `alfcy=0`, the radial mesh in the SOL is almost uniform except near the x-point; for details, see the user-callable function `rho1`.

For double-null configurations the radial distribution of flux surfaces can be different for the inboard and outboard legs of the plasma. In this case `psi0max` and `alfcy` refer to the outboard leg and one must supply additional input for

```

psi0max_inner :flux at 'outer' wall boundary on inboard leg

```

`alfcy_inner` :cluster factor for flux surfaces on inboard leg

For `kymesh=0` the user manually specifies the poloidal magnetic flux for each surface in the mesh; this is done by setting the dimensions, allocating and filling the arrays `psitop` and `psibot`; these must be consistent with `nycore`, `nysol`, `psi0min1`, `psi0min2`, and `psi0max` (and `psi0max_inner` if `geometry="dnbot"`). The array `psitop` contains the un-normalized values of the poloidal magnetic flux along a cut at the top of the mesh (or midplane if `geometry="dnbot"`); `psibot` contains the values along the lower divertor plates, starting at the outermost flux surface (nearest the centerpost) on the inboard half of the mesh and proceeding radially to the outermost flux surface on the outboard half of the mesh.

The poloidal distribution of the mesh is controlled by the following input parameters. The mesh ends at the divertor plates which are defined by the separatrix strike points included within in the input file `aeqdsk`. The mesh is divided into various regions, and one can specify the number of cells in each as follows:

<code>nxleg(1,1)</code>	:number of poloidal cells in inboard leg between divertor plate and x-point
<code>nxcore(1,1)</code>	:number of poloidal cells in inboard region of core
	:between x-point and top (or midplane for double-null)
<code>nxcore(1,2)</code>	:number of poloidal cells in outboard region of core
	:between x-point and top (or midplane for double-null)
<code>nxleg(1,2)</code>	:number of poloidal cells in outboard leg between divertor plate and x-point

Additional control over the distribution of the mesh along the separatrix in the poloidal or x-direction is provided by a function `x(t)` where `t` is the indexing parameter that labels the cells; the grid points are spaced uniform in `t` for a given region (leg, or core), and the actual spacing in `x` is determined by `x(t)`. There are a number of options for `x(t)` controlled by the switch `kxmesh`:

<code>kxmesh=0</code>	:manual definition of seed points
<code>kxmesh=1</code>	:use linear*rational form for <code>x(t)</code> in divertor
<code>kxmesh=2</code>	:use linear*exponential form for <code>x(t)</code> in divertor

kxmesh=3 :use spline form for $x(t)$ everywhere
 kxmesh=4 :use exponential+spline form for $x(t)$ in divertor

For kxmesh=0, one must fill the arrays seedxp and seedxpxl. These arrays give the location on the separatrix of the mesh points in percentage of the distance from the x-point to plate or x-point to top of machine, etc; thus, all values must lie between 0-100 and be monotonic. A standard procedure for setting the arrays seedxp and seedxpxl is to generate an approximate mesh with a kxmesh.ne.0 option, and then read the files rdgen.seedxp in directory uedge/in. This fills the arrays with the current mesh values, and one can then edit these arrays by inserting or moving points. The resulting arrays should be saved into a PFB file which can then be read in using the restore command for generating the desired mesh with kxmesh=0.

For kxmesh=1 or kxmesh=2 the poloidal spacing of the cells between the x-points is nearly constant unless one makes the variable slpxt larger than unity (1.2 to 1.3 is typical); having slpxt > 1 causes the cells to cluster near the x-point on both sides, and thus often match more smoothly with the cells in the divertor leg regions.

For kxmesh=4 the user specifies sub-regions in front of each divertor plate: the variables nxgas(1:2) specify the number of cells in the region, the variables dxgas(1:2) specify the size of the first cell at the divertor plate, and the variables alfx(1:2) specify the exponential factor for the cell-to-cell variation in the region. The relation between dxgas and the total length of the exponential region, L, is given by $L = dxgas * [\exp(alfx * nxgas) - 1] / [\exp(alfx) - 1]$.

4.2. *Non-orthogonal grids*

Non-orthogonal grids can be generated by setting the switch ismmon:

ismmon=0 :strictly orthogonal mesh (default)
 ismmon=1 :poloidal mesh is compressed/expanded w.r.t. orthogonal
 ismmon=2 :poloidal mesh varies smoothly on each flux surface
 ismmon=3 :combination of ismmon=0,1,2

This switch affects the distribution of meshpoints on each flux surface, but does not alter the the number or spacing of the flux surfaces.

For `ismmon=1` the mesh is generated by deforming a previously generated orthogonal grid along flux surfaces that intersect the divertor plates. The original mesh is uniformly compressed or expanded in the poloidal direction until the end of the mesh just coincides with the divertor plate. The compression or expansion occurs along each flux surface between some upstream reference surface (such as the midplane) and the divertor plate. A smoothing procedure may subsequently be applied to remove abrupt distortions in the mesh.

Options under the `ismmon=1` setting are:

<code>istream</code>	:definition of the upstream reference surface :=0 (default) -> midplane in SOL, cut in p.f. region :=1 user-defined
<code>iplate</code>	:definition of the divertor plate surface :=0 (default) -> orthogonal plate :=1 user-defined
<code>nsmooth</code>	:number of smoothing passes applied to each surface :=0 -> no additional mesh modification :=2 (default) recommended

The user defines the divertor plates via the arrays

`rplate1(1:nplate1)` `zplate1(1:nplate1)`

for the inboard leg of the divertor, and

`rplate2(1:nplate2)` `zplate2(1:nplate2)`

for the outboard leg of the divertor. Usually, one would read this information from a text file prepared specifically for the device being modelled. Examples of such files for DIII-D, CMOD, TPX, and ITER are available from the authors. NOTE: For complicated divertor

geometries it may be necessary to simplify the divertor plate definition to avoid intersecting any flux surface more than once. The current version assumes there is only one intersection.

The user defines the fixed upstream reference surface via the arrays

```
rupstream1(1:nupstream1)    zupstream1(1:nupstream1)
```

for the inboard half of the mesh, and

```
rupstream2(1:nupstream2)    zupstream2(1:nupstream2)
```

for the outboard half of the mesh. Usually, one would read this information from a previously prepared text file. For example, on open SOL flux surfaces one might choose to modify the mesh only downstream from the midplane and on private flux surfaces only downstream from the “cut” under the x-point. The file that does this is **upstream.mpc** available from the authors. The mesh in the core region would not be modified.

EXAMPLE: In addition to parameters for an orthogonal mesh, set the following:

```
ismmon=1      # switch on mesh modification
istream=0     # use default upstream reference surface
iplate=1      # flag indicates user will supply plate definition
read plate._device# define divertor plate surfaces
nsmooth=2     # use default smoothing of “radial” surfaces
```

For `ismmon=2` the normalized poloidal distribution of mesh points is the same on each flux surface. The distribution on the separatrix flux surface is defined according to the input parameter ‘`kxmesh`’ described earlier in this writeup. This poloidal distribution is then normalized in terms of the total poloidal connection length from the divertor plate surface to the top of the mesh (for open SOL flux surfaces) or the cut under the x-point (for private flux surfaces). Then, on each flux surface the poloidal distribution of mesh points is obtained by scaling the normalized separatrix distribution with the poloidal connection length for that surface. The resultant mesh is non-orthogonal even for orthogonal divertor plates.

Options under the `ismmon=2` setting are:

<code>istream</code>	:definition of the upstream reference surface :=0 (default) -> top of the mesh in SOL, cut in private flux region :=1 user-defined
<code>iplate</code>	:definition of the divertor plate surface :=0 (default) -> orthogonal plate :=1 user-defined
<code>nsmooth</code>	:number of smoothing passes applied to each “radial” surface :=0 -> no additional mesh modification :=2 (default) recommended

EXAMPLE:

In addition to parameters for an orthogonal mesh, set the following:

```
ismmon=1      # switch on mesh modification
istream=0     # use default upstream reference surface
iplate=1      # flag indicates user will supply plate definition
read plate._device_ # define divertor plate surfaces
nsmooth=2     # use default smoothing of “radial” surfaces
```

Finally, it is noted that application of certain boundary conditions can become inaccurate when the mesh is at a substantial angle to the radial wall. Consequently, it is recommended that the user activate a switch that reverts to orthogonal radial differencing between $i_y=n_y$ and $i_y=n_y+1$ mesh points, as well as $i_y=0$ and $i_y=1$ mesh points. The setting to activate this orthogonal boundary differencing is `isybdryog=1`. This boundary switch is especially effective in eliminating the possibility non-physical ion fluxes away from material walls.

4.3. Adaptive-mesh capability

The mesh can be modified in response to the plasma state so as to obtain better resolution in spatial regions where physics variables are changing most rapidly. At present, this capability is limited to a poloidal re-distribution of mesh points along each flux surface. The number and position of the flux surfaces is not changed, i.e. the “radial” resolution is fixed. The basic idea is to poloidally refine the mesh near a “flamefront” surface between the x-point and the divertor plate(s). This process is not yet automated so the user must manually perform certain steps to change the mesh and then obtain a plasma solution on the modified mesh.

The user-callable subroutine `meshff(region)` modifies a reference mesh stored in arrays (`cmeshx3`, `cmeshy3`) and writes the modified mesh into the arrays (`cmeshx`, `cmeshy`). Here, `region=1` for the inboard leg and `region=2` for the outboard leg. The mesh is modified only between the x-point and the divertor plate(s); the core and adjacent SOL regions of the mesh are unchanged. It is the user’s responsibility to store the appropriate data in (`cmeshx3`, `cmeshy3`) before calling `meshff`. After `meshff` completes, it is necessary to call subroutine `writeue` which converts the (`cmeshx`, `cmeshy`) data into (`rm`, `zm`) data and writes the file `gridue` that is read by the plasma package when `gengrid=0`.

The flamefront surface is defined by the user via the arrays `rff1(1:nff1)` and `zff1(1:nff1)` where (`rff1`, `zff1`) are the (R,Z) coordinates [m] of the `nff1` data points. The ‘1’ here refers to the inboard divertor leg; there are corresponding variables with ‘1’→‘2’ for the outboard divertor leg. Storage for the arrays `rff1` and `zff1` is dynamically allocated by setting `nff1` and then calling `gchange(“grd.Mmod”,0)` from the parser.

Input data for the flamefront (FF) mesh modification includes -

<code>isxtform</code>	:a flag for choosing one of three possible forms for :the distribution of mesh points along a flux surface
<code>iswtform</code>	:flag for combining the original and FF meshes with constant :weight factor (<code>iswtform=0</code>) or index-dependent weight factor (<code>iswtform=1</code>)
<code>cwtff</code>	:shape factor for the <code>iswtform=1</code> option

:on combining original and FF meshes.

For the inboard leg:

nff1,rff1(),zff1()	:number of data points and (R,Z) coordinates of FF.
slpxff1	:slope reduction factor for x(ix) at FF position :on each flux surface.
slpxffu1	:slope reduction factor for x(ix) at position upstream of FF on each flux surface.
slpxffd1	:slope reduction factor for x(ix) at position :downstream of FF on each flux surface.
nxdff1	:number of cells between FF and divertor plate on each flux surface.
wtff1	:maximum weight factor for combining original and FF meshes.

For the outboard leg: replace 1 by 2 in the variable names above.

The input data controls the form of the meshpoint distribution along each flux surface by specifying various shape factors for an analytic function $x(t)$ that gives the poloidal distance from the x-point as a function of the poloidal meshpoint index. Let (t_1, x_1) represent the x-point, (t_2, x_2) the flamefront, and (t_3, x_3) the divertor plate.

For isxtform=1 we use a piece-wise functional form; on $t_1 < t < t_2$ use the rational function:

$$x(t) = x_1 + (x_2 - x_1) * (t - t_1) / ((t_2 - t_1) + \alpha * (t_2 - t))$$

and on $t_2 < t < t_3$ use the rational function:

$$x(t) = x_2 + (x_3 - x_2) * (t - t_2) / ((t_3 - t_2) + \beta * (t_3 - t))$$

where α and β are chosen to give a specified slope at t_2 . The slope is expressed as the product of the average slope and a slope reduction factor $slpxff$, $x'(t_2) = slpxff * (x_3 - x_1) / (t_3 - t_1)$ where a '1' or '2' should be appended to 'slpxff' for the appropriate divertor leg.

The `isxtform=2` option uses a slightly more general form for $x(t)$ which allows the user to also specify the slope at the upstream point $(t1, x1)$ in the form $x'(t1) = \text{slpxffu} * (x3 - x1) / (t3 - t1)$ where a '1' or '2' should be appended to 'slpxffu' for the appropriate divertor leg.

The `isxtform=3` option uses a similar form for $x(t)$ which allows the user to specify the slope at all three data points via the slope factors `slpxff`, `slpxffu` and `slpxffd`:

$$x'(t1) = \text{slpxffu} * (x3 - x1) / (t3 - t1)$$

$$x'(t2) = \text{slpxff} * (x3 - x1) / (t3 - t1)$$

$$x'(t3) = \text{slpxffd} * (x3 - x1) / (t3 - t1)$$

where a '1' or '2' should be appended to 'slpxffu', 'slpxff' and 'slpxffd' for the appropriate divertor leg.

To facilitate a gradual transition from the original mesh to a flamefront modified mesh, the two meshes are combined via a weight function, $wt(t)$, to produce the final form of the meshpoint distribution $x(t)$:

$$x(t) = wt(t) * xFF(t) + (1 - wt(t)) * x0(t)$$

where $x0(t)$ represents the original mesh and $xFF(t)$ represents the flamefront mesh defined by the `isxtform` options above. The form of the weight factor is controlled by the flag `iswtform`: `iswtform=0` -> constant $wt(t) = wtff$ and `iswtform=1` -> smooth increase from 0 at x -point to $wtff$ at flamefront, where '1' or '2' should be appended to the 'wtff' for inboard or outboard regions.

The input parameter `nxdff` controls the number of cells downstream from the flamefront on each flux surface. At present, this number is the same for all flux surfaces. If `nxdff=0`, the number of downstream cells for the flamefront mesh is set equal to the number of downstream cells on the separatrix flux surface of the original mesh; otherwise, the user-specified value of `nxdff` sets the number of downstream cells on the flamefront mesh.

In summary, the steps necessary to use the adaptive mesh facility are:

1. define the reference mesh via the `grd` package arrays (`cmeshx3`, `cmeshy3`).

2. define the flamefront surface for each divertor leg.
3. set various flamefront mesh control parameters.
4. call subroutine meshff(region).
5. call writeue to convert the (cmeshx,cmeshy) data to (rm,zm).
6. set gengrid=0 (or mhdgeo=0 for cartesian configurations) and isnonog=1 before executing with `exmain`.

4.4. *Adding poloidal cells near the x-point*

The user may add extra poloidal cells near the x-point by setting the variables `nxxpt` and `nxmod`. Here `nxxpt` is the number of extra cells added between the x-point and the poloidal face `nxmod` indices away from the x-point for each of the four quadrants of a single-null divertor. This then results in a total of $4*nxxpt$ extra poloidal cells given by

$$\text{total poloidal cells} = nxleg(1,1) + nxleg(1,2) + nxcore(1,1) + nxcore(1,2) + 4*nxxpt$$

`nxmod` should be 1 or greater, with `nxmod=2` recommended; if `nxmod=2`, the two cells poloidally adjacent to the x-point are recalculated including the number of extra mesh points, `nxxpt`. Note that these cells are not strictly orthogonal, so it is safest to use the nonorthogonal difference stencil (`isnonog=1`), but the error in not doing so may be small and is limited to the modified x-point region.

There is some control over the spacing of these extra cells through the variables `alfxpt` and `alfxpt2`; `alfxpt` controls the nonuniformity of the poloidal mesh in the modified region, and `alfxpt2` controls how rapidly the poloidal face shape returns to a smooth arc as one moves away from the x-point. The practical range of `alfxpt` is roughly $0.25 < alfxpt < 1$, where `alfxpt=1` gives uniform spacing (the default) and `alfxpt=0.5` gives the cell faces closer to the x-point by roughly 0.707. The range of `alfxpt2` is $1 < alfxpt2 < 2$, where higher values force the mesh to return to a smooth arc faster. It is best to try a few values and then look at the result by plotting the mesh with the `plotmesh` script.

4.5. *Top-of-mesh/limiter option*

By default, the spatial extent of the “inboard” and “outboard” regions of the core/SOL are delimited by a vertical line from the magnetic axis upward through the separatrix. This top-of-mesh/limiter position can be changed by setting the switch `islimon=1` and defining the poloidal angle of the new top-of-mesh/limiter position, `theta_lim`. `theta_lim` should lie with the limits $-\pi < \text{theta_lim} < \pi$ and should not be too close to the x-point position. The outboard midplane position corresponds to `theta_lim=0` and the default is `theta_lim=pi/2`. The `islimon=1` option automatically turns on a procedure for checking the angular position of every data point on every flux surface, so the `flx` package may run slower.

4.6. *Cartesian and cylindrical configurations*

In addition to generating a mesh obtained from an MHD equilibrium code or model, UEDGE has the ability to simulate simpler configurations, namely, either Cartesian or cylindrical geometries. These options are controlled by the input variable `mhdgeo`; `mhdgeo=-1` is used for a Cartesian configuration and `mhdgeo=0` is used for a cylindrical configuration.

For the Cartesian case, one can set the following mesh parameters:

<code>radx</code>	:position outer "radial" wall, across-B-field direction
<code>radm</code>	:position of inner "radial" wall
<code>rad0</code>	:position of separatrix
<code>alfyt</code>	:radial nonuniformity factor; $< 0 \rightarrow$ expanding
<code>isfixlb</code>	:set left poloidal boundary as sym. plane
<code>za0</code>	:“poloidal” symmetry plane location
<code>zaxpt</code>	:“poloidal” location of x-point
<code>zax</code>	:“poloidal” location of divertor plate
<code>alfxt</code>	:poloidal mesh nonuniformity factor
<code>btfix</code>	:constant total B-field
<code>bpolfix</code>	:constant poloidal B-field

Within UEDGE, the variable `rm` gives the “radial” distance across the B-field and `zm` gives the poloidal distance.

For the cylindrical case, an annulus is simulated with a minimum radius of `radm` and a maximum radius of `radx`; the radial coordinate is `rm`. The axial distances are controlled by `za0` and `zax`.

The poloidal and toroidal components of the magnetic field for the Cartesian and cylindrical geometry are set according to

$$B_{pol} = B_{pol,fix}(r/R_{maj,fix})^{\sigma_{bpol}}, \quad B_{tor} = \sqrt{(B_{t,fix}^2 - B_{pol,fix}^2)}(r/R_{maj,fix})^{\sigma_{btor}},$$

using the user-given parameters `btfix`, `bpolfix`, `rmajfix`, `sigma_bpole`, `sigma_btore`.

5. Running UEDGE in the Time-Dependent Mode

UEDGE can use a couple of automated ODE integrators. The variable name that selects the integrator is called `svrpkg`, and one sets it by typing `svrpkg="iname"` where `iname` is one of the following:

```
iname = vodpk          # preconditioned Krylov package for ODE's
iname = daspk          # preconditioned Krylov package for ODE's & algebraic eqns
```

The default is `svrpkg="vodpk"`. This ODE solvers have been developed by G. Byrne [27] and A. Hindmarsh, and is available through the NETLIB web site <http://www.netlib.org/>.

There are also two Newton solvers with the names `nksol` and `newton`. More details of the Newton solvers are described following the next section on time-dependent simulations. However, we have found that using the `nksol` option with a timestep `dtreal` as described below in Sec. 5.5 is most robust. The NKSOL modified Newton solver (without the `dtreal` timestep) has been developed by P.N. Brown and colleagues at LLNL and follows the procedures given in Ref. 28. More recent replacements for VODPK and NKSOL which run on parallel computers are PVODE and KINSOL [29]. A version of UEDGE does run on parallel computers [30], but the operational details are not included in this manual.

It is very effective to precondition the Jacobian matrix for both the time-dependent and the steady-state Newton iterations. The options are discussed in the section on the Newton solver `nksol` below (search for the words `nksol` and `premeth`). Here we just mention that three options are available for the time-dependent mode, `premeth="banded"`, `"inel"`, or `"ilut"`.

5.1. *Setting simulation time and diagnostic output*

Output data concerning the performance of the time integration is stored in a sequence of evenly spaced logarithmic time intervals. The number of outputs is set by `isteps(1)`, which is defaulted to 100. The total simulation time is given by `trange*runtim` [sec]. The variable `runtim` gives the time increment for the first interval; `trange` and `runtim` are defaulted to `1.e+7` and `1.e-7`, respectively; the default total simulation time is thus 1.0 sec. Specifically, the output time corresponding to the cumulative output index, say `iout`, is `tout = (1.17489756)**iout * runtim`. The plasma variables are also stored at these output times in the arrays `nist1`, `upst1`, `test1`, `tist1`, `ngst1`, and `phist1` (see subroutine `uedriv` in file `odesolve.m`).

The most important thing to know from the last paragraph is that for the default settings, the total simulation time is `1.e7*runtim` seconds.

5.2. *Calculation of Jacobian*

We use the same subroutine, `pandf`, to evaluate the full right-hand sides of the ODE over the whole grid, and to calculate the Jacobian. In calculating the Jacobian, the range of the do-loops over the grid is restricted to the vicinity of the variable that is being perturbed. This range is controlled by three variables: `xlinc(=1)` is the incremental range to smaller `ix`, `xrinc(=2)` is the incremental range to larger `ix`, and `yinc(=1)` is the incremental range to both smaller and larger `iy`. One can test the Jacobian calculation by setting these to values larger than `nx+1` and `ny+1` so that the whole range of the do-loops is done for every perturbation; this is very inefficient for normal use, however.

5.3. Accuracy

The relative accuracy of the time-dependent integration is set by `rtolv`, which is a vector of length 30 to allow for the possibility of setting up a maximum of 30 different sequential runs where one might change `rtolv` and `runtim`, for example; in practice, we may use 2 or 3 for grid sequencing. For a given run `rtolv` is used to set the `vodpk` relative error variable `rtol`. We then define the `vodpk` absolute error variable `atol(i)=catol*rtol*(guess at solution for variable i)`. Here `catol` stands for a set of scale factors for each variable set, i.e., `cniatol`, `cupatol`, `cteitol`, `ctiatol`, `cngatol`, and `cphiatol`. Typically, we choose `rtolv`=1.e-3 or 1.e-4, which is large by usual `vodpk` standards. However, we want to reach steady state as quickly as possible. The usage of the large `rtolv` required us to add a variable to the Krylov solvers for calculating the vector $A*v$ by finite difference. The perturbation previously used for the finite difference was `rtol` on the assumption that the user would choose `rtol` $\sim \sqrt{\text{machine roundoff}}$ or $\sim 1.e-7$ for the Cray. As we may have `rtol` $\sim 1.e-3$, we let the perturbation in the Krylov solver be `srtolpk*rtol`, with the default `srtolpk`=1.e-4 used to give a perturbation of $\sim 1.e-7$.

5.4. Boundary conditions for the time-dependent mode

The boundary conditions are set in subroutine `bouncon` (see file `boundary.m`). For the solver `vodpk`, the boundary conditions are implemented as ODE's. For example, if we want the density `ni` to be `nb`, then the equation is

$$\frac{\partial n_i}{\partial t} = -cnurn \times nurlx(ni - nb) \quad (1)$$

where `nurlx`=1.e8 [sec] is a large relaxation frequency to force the boundary condition to be satisfied on a time scale short compared to the evolution of the non-boundary variables. The scale factor `cnurn` (=1 for default) applies only to the density equations. A similar equation is used to specify a flux-like boundary condition. The other variables have the same type of boundary equations and use the scale factors `cnuru`, `cnure`, `cnuri`, `cnurg`, and `cnurp` to allow independent adjustment for `up`, `te`, `ti`, `ng`, and `phi`, respectively. It should be noted that the potential equation, arising from $\nabla \cdot \mathbf{J} = 0$, has no time derivative when inertial and finite charge effects are ignored, and thus is treated this same way.

If one uses `daspk` as the solver, the boundary conditions are specified as algebraic equations. This guarantees that the boundary conditions are satisfied at each timestep and should be viewed as the preferable method. At the initial time, the algebraic equations must be satisfied to a given level of accuracy. This is now performed by effectively using the **NKSOL** Newton solver to satisfy only the boundary conditions and the potential equation if it is switched on (`isphion=1`). Note that if one wishes to temporarily freeze the potential equation and later turn it back on, set the variable `isphiofft=1` together with `isphion=0`. Then, to turn the potential evolution back on, reverse those two settings.

5.5. *Using the NKSOL solver in a time-dependent mode*

5.5.1. *Interactive mode*

It is possible to use `svrpkg="nksol"` in the time-dependent mode by setting `dtreal` to the desired timestep in seconds. Here a term is added to each of the non-boundary equations to account for a linear, or backward Euler, time advance; i.e., $d(y_l)/dt \rightarrow (y_{l_new} - y_{l_old})/dtreal$. It is possible to also add the `dtreal` term to the boundary equations and the potential equation by setting the flag `isbcwdt=1`; this can be useful for relaxing the complete system far from equilibrium and is similar to what `vodpk` does. The number of such timesteps is controlled by the parameter `nsteps_nk` (defaulted to 1). If `nsteps_nk > nsteps`, the time-dependent arrays `test1(istep,ix,iy)` will be filled until the number of steps exceeds `nsteps`, and then the last value will be overwritten with the last value. This option is similar to the pseudo timestep method for **NKSOL** described below which is controlled by the parameter `dtnewt`. Note that these cannot be used together, and an error message is issued if `dtreal` and `dtnewt` are simultaneously less than 1.e5 seconds.

5.5.2. *Controlling NKSOL time-dependence with scripts*

Two script files, called `rdinitdt` and `rdcontdt`, are used for running `uedge` in a time-dependent mode. The way to use these is as follows:

1. Initialize a few new variables for the time-dependent mode by typing `read rdinitdt`.
2. Set the initial timestep called `dtreal`; usually `dtreal=1e-9` is conservative.
3. Run `uedge` to a converged state by typing `exmain`.
4. If the run converges, begin the time-dependent run by typing `read rdcontdt`.

Some more details: the time-dependent mode automatically runs a series of `uedge` problems where the timestep `dtreal` has been added to the equations, so each step corresponds to a real timestep. The script `rdcontdt` contains two main loops, one inner loops keeps `dtreal` fixed for `ii2max` iterations (default=5), and the outer loop increases `dtreal` by a factor `mult_dt` (default=3.4). If a given iteration fails, the script automatically reduces the timestep by `mult_dt` and tries again. If you see `dtreal` getting reduced below $1e-10$, this whole procedure will likely fail, and it is generally recommended that you kill the job. The run stops when either the total accumulated time, `dt_tot`, is equal to `t_stop` (default=10 s, which is almost always very close to a steady state), or when the number of outer loop iterations exceeds `ii1max` (default=100).

At each timestep, the solution is saved to the pfb file called `pfdt_savefname`, where `savefname` is a user specified 5-place character string initially set in `rdinitdt` to `savefname="it333"`. After reading `rdinitdt`, you can reset `savefname` to any 5-character string you want. Note that this pfb file is continually overwritten, so if you have one at the end of a run that you want to save for future restarts, it is safest to change its name.

You can also save a time-history of the run by setting the variable `n_stor` in `rdinitdt` to a non-zero value; typically, `n_stor=100` is used to get a first look if the run has interesting time-dependence. The users also controls the time window where the data is stored by setting `tstor_s` and `tstor_e`, the beginning and ending time of the time window. The script begins storing the solution when the `dt_tot = tstor_s`, and stores as close as it can for the time increments $(tstor_e - tstor_s)/(n_stor - 1)$. The actual time corresponding to each data storage is recorded in the array `tim_stor(1:n_stor)`. The data that is stored is the primary variables `ni`, `up`, `te`, `ti`, `ng`, and `phi` in the following arrays:

```

ni_stor(1:n_stor,0:nx+1,0:ny+1,1:nisp)
up_stor(1:n_stor,0:nx+1,0:ny+1,1:nisp)
te_stor(1:n_stor,0:nx+1,0:ny+1)
ti_stor(1:n_stor,0:nx+1,0:ny+1)
ng_stor(1:n_stor,0:nx+1,0:ny+1,1:ngsp)
phi_stor(1:n_stor,0:nx+1,0:ny+1)

```

This data is written to another pfb file called `pftstor.savefname`, so you can save it for latter viewing; be careful to change the name if you really want to save it, since it will get overwritten by the next run unless you change the 5-character name `savefname`.

6. Running UEDGE with a direct Newton iteration to steady state

6.1. Switches and diagnostic output

To invoke the simple direct (non-Krylov) Newton iteration, set `svrpkg="newton"` (note that this option is now rarely used). The code then uses the subroutine `newton` to update the solution in the form

$$\delta y_l = -J^{-1}F(y_{l_{old}}) \quad (2)$$

where F is the right-hand side of the ODE's, δy_l is the change in y_l , and J is the Jacobian. At each iteration, the code prints out two lines of diagnostics. Here, `sumnew1` is the average change in the magnitude of the y_l 's for this iteration, `sumr1dy` is the average of $\text{abs}(\delta y_l/y_l)$, and `saux2` is the fraction of the Newton update allowed based on the variable `rlx`. Here `rlx` is the maximum amount that any y_l is allowed to change relative to its old value for a given iteration. The default is `rlx=0.4`. The second output line gives `sumf`, the average of the magnitude of the right-hand-sides, `ivmxchn` is the `ieq` index of the $y_l(\text{ieq})$ that has the largest magnitude of $\delta y_l/y_l$, and the `(ix,iy)` gives the location on the grid for this y_l variable.

6.2. Preconditioning for the `svrpkg="newton"` case

Only `premeth="banded"` works for this option. It uses the direct banded solver SG-BFA. An error message will be received if any other option is specified for `premeth` when `svrpkg="newton"`.

6.3. Determining convergence trends and abort command

The Newton iteration should show a clear trend toward convergence, i.e., `sumnew1` and `sumf` decreasing after 5 to 10 iterations. If this does not occur, experience shows that convergence is very unlikely. To abort the iteration, type `ctrl-c` which will return you to the `DEBUG>` prompt; following the instructions that appear on the terminal, you may interrogate `UEDGE` and then type `cont` to continue, or type `abort` to return to the `UEDGE>` prompt. Another good measure of convergence is the initial value of `saux2`; if `saux2 < 1.e-2`, convergence is very unlikely, if $1.e-2 < \text{saux2} < 1.e-1$, convergence is somewhat likely, and if $\text{saux2} > 1.e-1$, convergence is quite likely.

The criterion for convergence is that $\text{sumnew1} < \text{rwmin}$, with `rwmin=1.e-11` as the default. Other useful control variables are `nmaxnewt` which is the maximum number of iterations that the code will try (with a absolute hardwired upper limit of 101). The variable `scrit` causes the old Jacobian to be used if the average of $\text{abs}(\delta y_l / y_l)$ is less than `scrit`; the default is `scrit=1.e-4`.

7. Running UEDGE with Krylov-Newton Iterations to Steady State

7.1. Switches and diagnostic output

To invoke the preconditioned Krylov Newton solver option based on Peter Brown's NKSOL package, set `svrpkg="nksol"`. The `nksol` option is generally preferred over the `newton` option. There are a variety of options for this package that are briefly described in the variable descriptor file `bbb.v` (groups `Lsode` and `Ilut`). Commonly changed flags are `mf nksol`,

mdif, and mmaxu:

mfnsol = 1	:dogleg search strategy using GMRES iterative solver
= 2	:linesearch with Arnoldi method iterative solver
= 3	:linesearch with GMRES method (default)
mdif = 0	:Matrix-vector multiply $J*v$ is approximated by numerical :finite difference of RHS (default)
= 1	:Matrix-vector multiply $J*v$ is done directly with current J
rlx = 0.4	:restricts relative change of density and temperatures to be :less than rlx at each point
stepmx = 1.e9	:restricts global change of variables, $\text{sum}[\text{sqrt}[(\text{del}(u)**2)]]$, :to be less than stepmx. Can be used instead of rlx.
itermx = 30	:Maximum number of nonlinear iterations
incpset = 5	:Maximum nonlinear iteration before Jacobian is reevaluated
mmaxu =	:Now calculated internally with the algorithm $\text{mmaxu}=\text{neq}**0.5$:To set a specific value at input, set $\text{ismmaxuc}=0$ (default is 1).
epscon1 = 0.1	:Use to define tolerance of linear iterative matrix solutions :with the algorithm $\text{epsfac}=\text{epscon1}*\min(\text{epscon2}+\text{frnm})$. Final :tolerance is $\text{epsfac}*\text{frnm}$
epscon2 = 1.e-2	:Use to define tolerance of linear iterative matrix solutions :with the algorithm $\text{epsfac}=\text{epscon1}*\min(\text{epscon2}+\text{frnm})$. Final :tolerance is $\text{epsfac}*\text{frnm}$

Note that negative values of `mfnsol` have the same meaning as the positive ones, except that the global constraints are not used (so `frnm` can increase substantially from one iteration to another).

7.2. Preconditioning options

An important part of the Newton iteration is the preconditioning of the Jacobian matrix. There are several methods available and these are controlled by the flag `premeth="iname"`;

the different options for the preconditioner are:

premeth= "banded"	:uses the direct banded solver SGBFA. Requires a lot of :storage for large problems, but is fast on the Cray for :moderate size problems
premeth= "inel"	:uses a partial LU decomposition with fill-in on existing :diagonals of Jacobian only; called ILU0
premeth= "ilut"	:uses a partial LU decomposition about existing elements of :Jacobian - not based on diagonals only. Amount of fill-in controlled :by lfililut; typical problems require lfililut=3 to 100.

The output data on performance of the `nksol` routine at each nonlinear iteration is controlled by the flag `iprint`. No output occurs if `iprint=0`. If `iprint=1` (default), the iteration count, norm of the residual, and number of right-hand-side evaluations are printed, indicating roughly the number of linear iterations. If `iprint=2`, detailed data concerning the linear iteration is printed for each nonlinear iteration: the norm of the residual and the value of norm required to meet convergence test. Also, if the constraint condition preventing negative densities or temperatures, or a relative step size is too big $[\text{Del}(u)/u > \text{rlx}]$, resulting in a reduced step size, the message `ivio=1, pnrn=...` will appear.

7.3. Row and column scaling and rescaling

Several techniques are used to improve the numerical properties of the Jacobian preconditioning matrix and to better condition the nonlinear Newton problem. The most straightforward is scaling the rows of the Jacobian by the largest element in the row; this is effected by the switch `issfon=1` (default) which generates the scaling vector `sfscal`. This scaling is actually applied to the Jacobian if `isrnorm=1` (default).

Column scaling is a more recent addition (mid-1995) and is turned on by the switch `iscolnorm` (default=0). Presently, either `iscolnorm = 2` or `3` is recommended, and both have nearly the same effect; `2` completely disregards the old global scaling and `3` does the column

scaling after the global scaling. The column scaling is essentially a local normalization of the variables over the mesh to be of order unity, which improves the numerical solve-ability of the system.

Rescaling of the Jacobian matrix is activated by the switch `ireorder=1` (default). This causes the Jacobian to be reordered using the reverse Cuthill-McKee algorithm. Typically it reduces the number of linear Krylov iterations by 30-50%, but can make the difference between convergence and no convergence.

7.4. *Pseudo-transient timestep*

A pseudo timestep can be added to the Jacobian for `svrpkg="nksol"` which generally increases the radius of convergence (one can take larger steps away from existing solutions), but decreases the rate of convergence. The timestep adds a diagonal term to the left-hand side of the linear matrix equation but not to the right-hand side. Specifically, consider the time-dependent equation for a vector of variables \mathbf{x} of the form $d\mathbf{x}/dt + \mathbf{f} = 0$. Performing a Taylor series expansion gives the Jacobian, \mathbf{J} , and

$$\frac{(x_n - x_o)}{dt} + \mathbf{J} \cdot \mathbf{x}_n = -f(x_o) \quad (3)$$

where x_o and x_n are the variable values at the old and new timestep, respectively. The pseudo transient technique neglects the $-x_o/dt$, but retains x_n/dt , yielding the equation

$$(\mathbf{I}/dt + \mathbf{J}) \cdot \mathbf{x}_n = -f(x_o) \quad (4)$$

where \mathbf{I} is the identity matrix. This additional term is added to both the preconditioning Jacobian, and to the $\mathbf{J} \cdot \mathbf{x}$ finite-difference Jacobian-vector product calculated in the Krylov algorithm.

In `UEDGE`, the pseudo timestep, dt , is called `dtnewt`. The default value for `dtnewt` is 1.e20 seconds which effectively removes the \mathbf{I}/dt term. To use this technique, it is typical to start with `dtnewt=1.e-5` to `1.e-4`, and run for about 10 iterations (set `itermx=10`). The residual as measured by `frn` should be decreasing, but do not expect convergence. Then update the “save” variables by doing a `read reset`, increase `dtnewt` by a factor of 3 or 10, and

repeat the procedure. By the time you get to $dt_{newt}=1.e-3$, or so, the convergence should be accelerating, and one can often increase dt_{newt} by larger factors; typically, $dt_{newt}=0.1-1.0$ is almost equivalent to $dt_{newt}=1.e20$. As of yet, there is no systematic procedure coded for automatically running through the sequence described above, although Knoll and McHugh have had some success with the Switched-Evolution Relaxation (SER) method.

8. Boundary Condition Options

8.1. *Specifying gas input and pumping on the side-walls*

One can specify up to 10 sources (or sinks - pumping regions) on the outer wall ($iy=ny$) and 10 on the inner wall ($iy=0$) by setting the variable `nwsor` to the number desired. The sources come in pairs, but the inner and outer parameters can be specified independently. These gas boundary conditions are set up in subroutine `walsor`. Each source uses a variable called `igspsoi(i)` or `igspsoo(i)` for the inner and outer wall sources, respectively that defines which gas species the source contributes to. Thus, `igspsoi(i)=j` means that source `i` contributes to gas species `j`.

The location of the sources is set by `xgasi(i)` and `xgaso(i)` for inner (private flux) and outer walls, respectively. To set these input parameters, it is helpful to examine the arrays `xfwi` and `xfwo` which give distances along these flux surfaces. If `issorlb(i)=1`, the distances `xgasi,o(i)` [m] are measured from the inner, or left divertor plate; if `issorlb(i)=0`, the distances are measured from the outer, or right divertor plate. The total width of the region is given by `wgasi,o(i)`; if `igasi,o(i) > 0`, the source is taken to be a gas source with a cosine shape over the defined region, going to zero at the edge. If there are multiple, finite strength sources present, the net source at a given location is the sum of all of the overlapping sources.

The special setting of `igasi,o(i)=0` is used to specify a pumping region with a uniform albedo defined by `albdso,i(i)` over the region of the source. The albedo of the side walls are specified for each gas species separately as follows: For each source which has `igasi,o(i)=0`, the second variable, `igspsoi,o(i)`, defines which gas species the source defines the albedo for,

just as for finite gas sources (see above). For example, if you want to set the albedo (for the private-flux wall) for gas species 1 to 0.95 and that of gas species 2 to 0.90, you should use two sources by setting `nwsor=2`, and

```

igasi(1:2) = 0
xgasi(1:2) = 0
wgasi(1:2) = 1000.
igspori(1) = 1
albdsi(1) = 0.95
igspori(2) = 2
albdsi(2) = 0.90

```

Here `wgasi` is set to a large number to span the whole simulation region. The gas input depends on the number of gas species used. If only one gas species is used (`ngsp=1`), the current `igasi,o(i)` is the boundary condition for that single species in the given region and one needs to account for simultaneous puffing and pumping by reducing the net gas input.

The wall sources can also be used to redistribute the gas flux absorbed at one location by injecting it at another location. Here the neutral flux is calculated from the albedo defined over a specified source region, and then reinjecting as a gas flux over a different (or the same) region with a cosine distribution characteristic of the sources. For this option, the gas flux in a given source region is the sum of the redirected flux from a “sensing” source (remote or local) and the albedo associated with the given source region. As an example, consider the current produced by the albedo `albdsi(k)` over the region specified by source `k` with location and width `xgasi(k)` and `wgasi(k)`, respectively. To reinject this current over the region specified by source `j`, with location and width `xgasi(j)` and `wgasi(j)`, set `ncpli(k)=j` to establish the coupling, and set `cplsori(j)` equal to the fraction of the current that will come through the source with index `j` [`cplsori(j)=1` gives the full current at source `j`]. Note that one may use `k=j`, and that the local albedo is included in determining the total flux.

It is now possible to measure the gas flux out of some region on the inner (outer) wall and reinject it on the outer (inner) wall. The only difference from the procedure described

in the last paragraph is that `ncpli` or `ncplo` are set to the negative index of the source. This signals the code that you want to measure the current on one wall over a region defined by, for example, `albdsi(i)` with location and width `xgasi(k)` and `wgasi(k)`, and inject it as source `j`, with location and width `xgaso(j)` and `wgaso(j)`. This connection is accomplished by setting `ncplo(k)=-j` and `cploro(j)=1` (for the full current) in the present example. This type of coupling could result in a degradation of the code performance since this inner/outer wall coupling is not included in the preconditioning Jacobian; however, a few simple tests suggest this may not be too serious of a problem.

One can also specify the side walls as material surfaces that emit recycled gas. This is done by switching on the flag `matwsi,o(i)=1` for a given wall source: the poloidal logic arrays `matwalli,o` are calculated internally based on the sum of the various `matwsi,o(i)` flags. The gas input at the boundary for `matwsi,o(i)=1` is determined by the wall recycling coefficients `recycw(1)` for hydrogen. As of 2001, the additional source from recycling is added to the local sputtering (see below in Sec. 8.6), albedo, and finite external sources from `igasi,o` fluxes. If `matwalli,o=0` in a given region, then only the finite external sources through `igasi,o` have any affect. The models for material side walls are still evolving, so users should check with the authors if they intend to use this option. The setting of the wall boundary conditions is often subject to user (and developer!) errors, so the user should always verify *a posteriori* that the gas fluxes on the boundaries are those sought.

The end plates are also sources of neutrals, where the gas flux is specified as `-recycp*(ion flux)`. Again, `recycp(1)` refers to the flux into the hydrogen gas, and `recycp(2...)` refers to impurity gases.

8.2. Other side-wall boundary condition options

There are several options to use for the density and temperature boundary conditions on the private-flux wall (`iy=0`) and the outer wall (`iy=ny+1`). These can be set through the switches `isnwconi,o` for the densities where the final `i,o` refer to the “inner” wall (`iy=0`) and the “outer” wall (`iy=ny+1`); note that the inner wall corresponds to the private-flux boundary for

divertor tokamaks. For the temperatures on these two wall surfaces, the switches are named `istewc`, `istiwc`, `istepfc`, and `istipfc`. Here “ti” and “te” refer to ion and electron temperatures, while “wc” and “pfc” refer to the outer and inner walls, respectively. First, we describe in some detail some of the more common settings for these switches, and then a short, cryptic list of all the options.

Constant value (Dirichlet) boundary conditions:

For ion density, set `isnwconi=1` for the private flux boundary (`iy=0`) and `isnwcono=1` for the outer wall boundary (`iy=ny+1`). The density values can be set through the poloidal arrays `nwalli` and `nwallo` for the inner and outer walls, respectively. Be careful if the mesh is increased and interpolation is used, as one must then update `nwalli` and `nwallo`.

For electron and ion temperature, set `istepfc=1` and `istipfc=1` for the private-flux boundary, and set `istewc=1` and `istiwc=1` for the outer wall boundary. Each of these switches can be set separately. The temperatures are then set to `tedge` (eV) or can be given separate and poloidally varying values through the arrays `tewalli`, `tiwalli`, `tewallo`, and `tiwallo`, where the final letter denotes inner(i) or private-flux boundary and outer(o), or outer wall boundary. Before you set `tewalli`, etc. to nonzero values, you must allocate memory for these arrays by typing `allocate` at the basis prompt. In using this option, be careful of interpolating to a larger grid, as the `tewalli`, etc will have to be manually interpolated after the `allocate`.

For the parallel velocity, one can specify zero value on the inner and outer walls by setting the switches `isupwi=0` and `isupwo=0`, respectively. The default value for each of these is unity, which gives the slip boundary condition, *i.e.*, zero radial derivative ($d u_p/dy = 0$).

Flux (Neumann) boundary conditions:

For ion density, set `isnwconi(i)` and/or `isnwcono(i)=0` and `ifluxni=1`. Note that index `i` corresponds to the difference ion density species.

For electron and ion temperature, set `istepfc=0`, `istipfc=0`, `istewc=0`, and `istiwc=0`. This results in the normal derivative being set to zero, $dT/dy=0$.

Extrapolation boundary conditions:

For the density and temperatures, it is possible to set extrapolation boundary conditions at $i_y=0$ and $i_y=n_y+1$. This condition sets the boundary value to be a linear extrapolation of the previous two points in the radial direction. These boundary conditions correspond to $isnwconi,o(i)=2$ and $istewc=2$, *etc.*

Short listing to settings:

For the inner and outer wall densities:

$isnwconi,o=0$, old case; if $ifluxni=0$, $dn/dy=0$; if $ifluxni=1$, $fniy=0$
 $isnwconi,o=1$, fixed density to $nwallo(ix)$ array
 $isnwconi,o=2$, extrapolation B.C.
 $isnwconi,o=3$, approx grad-length $lyni$, but limited by $nwimin$ and $nwomin$

For the outer wall temperatures (te and ti):

$iste,iwc=0$, set zero energy flux
 $iste,iwc=1$, set fixed temp to $tedge$ or $te,iwallo$
 $iste,iwc=2$, use extrapolation BC
 $iste,iwc=3$, set Te,i scale length to $lyte$ or $lyti$
 $iste,iwc=4$, set $feey = bceew*fniy*te$ or $feiy=bceiw*fniy*ti$

For the inner or private-flux wall temperatures, the switches $iste,ipfc$ the settings have the same meaning as for $iste,iwc$.

8.3. End-plate boundary condition options

The ion parallel velocity is taken to be the sound speed multiplied by the user-set scale factors $csfacib$ and $csfacrb$ at the inner and outer divertor plates, respectively. If the switch $isupss=1$, then the parallel velocity is allowed to be supersonic at the plate if the solution seeks this state.

The recycling coefficients at the plates can be made a function of radial position or an

albedo can be specified over a limited region to pump gas through the plate. The following variables are used to set the recycling and albedos:

<code>ndatlb(igsp,1)</code>	<code># number of data points along inner plate; if=0, recycling</code> <code># is uniform and specified by recycp*recycfi</code>
<code>ndatrb(igsp,1)</code>	<code># number of data points along outer plate; if=0,</code> <code># recycling uniform and specified by recycp*recycfo</code>
<code>ydatlb(igsp,idat,1)</code>	<code># dis. from inner sep. of data point for rdati & albpi</code>
<code>rdatlb(igsp,idat,1)</code>	<code># value of recycling coeff. at inner plate ydati;</code> <code># in between data points, linear interp. is used</code>
<code>adatlb(igsp,idat,1)</code>	<code># value of albedo at data point ydatlb; lin. interp used</code>
<code>ydatrb(igsp,idat,1)</code>	<code># outer-plate counterpart to ydatlb</code>
<code>rdatrb(igsp,idat,1)</code>	<code># outer-plate counterpart to rdatlb</code>
<code>adatrb(igsp,idat,1)</code>	<code># outer-plate counterpart to adatlb</code>

If the albedo in any segment along the plate is less than unity (the default), then the albedo boundary condition for the gas takes precedent over the recycling boundary condition. Note that it only makes sense to use at least two data points on the inside or outside because linear interpolation is used between the points. Operationally, one needs to first generate the mesh and run through `nphygeo` (just do a very short calculation) to generate the mesh locations relative to the separatrix on the plates; these are `yylb` and `yyrb` for the inner plate (left boundary) and outer plate (right boundary), respectively. With this information, you can decide where to put your data points (`ydati` and `ydato`).

The energy equations, including the drift energy, have plate boundary conditions set by the energy transmission factor `bcee` and `bcei` for electrons and ions, respectively in the simplest cases with `newbcl=0` and `isphion=0`. Here the energy flux is set to $feex = bcee \cdot te \cdot fnex$ and $feix = bcei \cdot ti \cdot fnix$, with $fnex$ standing for $ne \cdot vex \cdot sx$. With `isphion=1` and `newbcl=1`, the sheath potential contribution to the electron energy transmission factor is computed from the solution using sheath condition consistent with the parallel current at the plates, and the electron kinetic energy part is $2 \cdot Te$ and the ion kinetic energy factor is $2.5 \cdot Ti$.

8.4. *Boundary conditions at the core interface*

The ion density boundary condition is controlled by the variable `isnicore(ifld)`, where `ifld` is the index of the ion density, `ni(ifld)`. Thus, the settings of `isnicore` correspond to:

- `isnicore=1`: set uniform, fixed density, `ncore`
- `isnicore=0`: set flux to `curcore`, `ni` const only if diffusive
- `isnicore=2`: set flux & `ni` over range
- `isnicore=3`: set integrated flux, `ni` const for all cases
- `isnicore=4`: use impur. source terms (impur only)

The gas density boundary condition is controlled by the variable `isngcore(igsp)`, where `igsp` is the index of the gas species, `ng(igsp)`. Note that inertial neutral gas is controlled by `isngcore(1)`. Thus, the settings of `isngcore` correspond to:

- `isngcore=0`: set zero flux
- `isngcore=1`: set uniform, fixed density, `ngcore`
- `isngcore=2`: set rad. grad. to $\sqrt{\text{lam}_i \cdot \text{lam}_{cx}}$
- `isngcore=3`: extrap. for diff. gas only
- `isngcore=anything else`: set zero deriv which was prev default for inertial hydrogen

In restarting from a `isnicore=0` case (zero flux), use `mfnksol=-3` if `svrpkg="nksol"`. For the core temperature boundary conditions on `Te` and `Ti`, one may set either a specified power for electrons and ions as `pcoree,i` in Watts and setting the switch `iflcore=1`. To use fixed temperature boundary conditions, set `iflcore=0`, and then `tcoree` and `tcorei` give the electron and ion temperatures on the boundary in eV, respectively. The parallel velocity boundary condition is set by `isupcore`, where

- `isupcore = 0`: set `up=0` on core edge
- `isupcore = 1`: set $d(\text{up})/dy=0$ on core edge
- `isupcore = 2`: set `uu=0` on core edge, where `uu` is poloidal velocity.

8.5. Core boundary conditions with cross-field drifts

UEDGE only uses the ∇B portion of the diamagnetic drift velocity in the "body" of the computational region since the divergence of the other portion is identically zero. However, on the boundary surfaces, one must resort to the full diamagnetic expression for obtain particle and power fluxes. The correction for the radial current is available in older UEDGE version, but requires proper setting of switches. The names of all the present switches, and their present default settings are as follows:

```
cfniybbo = 0 # factor to includ. vycb in fniy,feiy at iy=0 only
cfniydbo = 0 # factor to includ. vycp in fniy,feiy at iy=0 only
cfeeybbo = 0 # factor to includ. vycb in feey at iy=0 only
cfeeydbo = 0 # factor to includ. vycp in feey at iy=0 only
cfqybbo = 0 # factor to includ. fqyb in core current B.C. only
cfqydbo = 0 # factor to includ. fqyd in core current B.C. only
```

These defaults settings are appropriate for cases without cross-field drifts. With cross-field drifts, the recommended settings are as follows:

```
cfniybbo = 0
cfniydbo = 1
cfeeybbo = 0
cfeeydbo = 1
cfqybbo = 0
cfqydbo = 1
```

8.6. Sputtering boundary conditions for the gas species

There can be wall or plate sources of gas that arise from sputtering, either physical or chemical. These are controlled by two flags for each gas species, `isph_sput(igsp)` for physical sputtering and `isch_sput(igsp)` for chemical sputtering.

8.6.1. *physical sputtering by ions on plates*

We consider two settings for the switch `isph_sput(igsp)`.

For `isph_sput(igsp) = 0`:

The sputtering of gas species `igsp` is controlled by the simple yield factor `sputtr` if the sputtering arrays `sputto(iy,igsp)` and `sputti(iy,igsp)` are zero as at the initialization of a run. Then `sputto,i ← sputtr`. However, to change the sputtering after it has been set initially by `sputtr`, you must directly change the arrays `sputti(iy,igsp)` for the inner plate and `sputto(iy,igsp)` for the outer plate.

For `isph_sput(igsp) = 1`:

This setting uses the DIVIMP/JET model obtained from David Elder. To use this properly, you must set the following input parameters:

<code>cion</code>	# atomic number of the target material; default is 6 for carbon
<code>cizb</code>	# max charge state of plasma ions; default is 1 for hydrogen
<code>crmb</code>	# mass of plasma ions in AMU; default is 2. for deuterium

The resulting yield along the divertor plate is put into the arrays `sputflxlb(iy,igsp,jx)` and `sputflxrb(iy,igsp,jx)`.

The number of ion species contributing to the sputtering is controlled by the input `ipsputt_i` and `ipsputt_e` denoting the beginning and ending indices of the ion density species doing the sputtering.

8.6.2. *chemical sputtering by neutrals on side walls*

Similarly, on the side wall, we use the switch `isch_sput`, but it now has more than just two options.

For `isch_sput(igsp) = 0`:

This setting allows the user to specify the chemical sputtering yield by initializing `chem-`

sputi,o(i,j), where the resulting flux boundary condition for gas species i is then

$$\text{fngy}(\text{igsp}=\text{i}) = \text{Sum_j} [\text{chemsputi,o}(\text{i,j}) * \text{ng}(\text{j}) * \text{vt} * \text{sy}]$$

For `isch_sput(igsp) > 0`:

Note that `isch_sput(igsp)` should be nonzero for only one gas species and this species should be carbon. These settings (1-7) use various models for the chemical sputtering of carbon from the side walls; this package comes from DIVIMP via David Elder (Univ. Toronto, private comm., 1998). The various models are:

<code>isch_sput</code>	Options for chemical sputtering:
1	Garcia-Rosales' formula (EPS94)
2	according to Pospieszczyk (EPS95)
3	Vietzke (in Phys. Processes of Interaction Fusion Plasma with Solids
4	Haasz (Submitted to J.Nucl.Mater., Dec. 1995)
5	Roth & Garcia-Rosales (Nucl. Fusion, March 1996)
6	Haasz 1997 (Brian Mech's PhD Thesis)
7	Haasz 1997 + reduced 1/5 from 10->5 eV (Porter)

It is recommended that `isch_sput = 5` or `6` be used, although recently G.D. Porter has a new fit (7) which departs from (6) at low energy, and is a better fit to the (Haasz) data at the low energy. The adjustment to the low-energy yield for `isch_sput = 7` is controlled by the parameter `redf_haas`, which has a default value of 0.2 Other input is the temperature of the surface, `t_surf`, in degrees K; the default is 300 K. The resulting chemical sputtering yield is stored in the arrays `yld_carbi,o(ix)` along the inner and outer walls.

8.6.3. *physical and chemical sputtering by ions on side walls*

Beginning with Version 4.31, there is the capability to include physical and chemical sputtering by ions on the private-flux and outer side walls. The input parameters are as follows:

<code>isi_sputw(igsp)</code>	Determines the outer wall sputtering for ions producing gas species <code>igsp</code>
<code>isi_sputpf(igsp)</code>	Determines the private flux wall sputtering for ions producing gas species <code>igsp</code>

Options (same for `isi_sputpf`):

<code>isi_sputw(igsp) = 0</code>	No ion sputtering (old case)
<code>isi_sputw(igsp) = 1</code>	Includes physical sputtering from ions
<code>isi_sputw(igsp) = 2</code>	Adds chemical sputtering from ions using model <code>isch_sput(igsp)</code> ; must set <code>isch_sput</code> to nonzero value (e.g., 7)

The sputtered fluxes multiplied, by cell-face area, from incident ions on the walls are computed and stored in the following two arrays:

<code>sputflxw(ix,igsp)</code>	units of particles/sec (positive)
<code>sputflxpf(ix,igsp)</code>	units of particles/sec (negative)

where `ix` is the usual poloidal index.

9. Sources and Sinks

It is possible to specify fixed particle, current, and energy sources having specific locations and Gaussian widths through the arrays `volpsor`, `voljcsor`, `pwrsoe`, and `pwrsoe`. These are part of the variable group `Volsrc`, and various control parameters are briefly described in the variable descriptor file, `bbb.v`.

The sources and sinks are normally determined by ionization of neutral gas and recombination of ion-electrons into neutrals. A special background source is used to prevent the neutral density from becoming too small. The gas continuity has the form

$$\frac{\partial n_g}{\partial t} + \nabla \cdot (n_g \mathbf{v}_g) = -n_{uiz}(n_g - bgsor) \quad (5)$$

where $n_g = ng$ is the gas density, and `bgsor` is given by `bgsor = ngbackg*(0.9 + 0.1*(ng-backg/ng)**ingb)`. Normally, `nbackg=1.e15 m**(-3)`, and `ingb=0` for defaults. However,

sometimes it is useful to set `ingb=2` or larger to prevent “pump out” of low density cells.

10. Flux-Limiting Transport Coefficients

10.1. Basic flux-limit expressions

The flux limits used in UEDGE for the parallel transport are of the form

$$\chi = \chi_s / [1 + |q_s/q_f|^{flgam}]^{1/flgam} \quad (6)$$

where χ_s is a classical (Spitzer) diffusion transport coefficient, and `flgam=1` is the default value (seldom changed). The second heat flux, q_f is the free-streaming flux defined by $q_f = flalfe * ne * v_{te} * te$ for electron thermal transport, where `flalfe` is a parameter often set to 0.21 to match some kinetic modeling, `ne` is the electron density, $v_{te} = \sqrt{te/me}$, and `te` is the electron temperature. If `flalfe=1e20` (the default), the flux limiting is effectively switched off. The ion heat flux is limited in the same manner with $e \rightarrow i$, and the ion parallel stress is limited to $flalfv * ni * ti$. The three flux limit factors are thus:

<code>flalfe</code> (recommend=0.21)	:for electron parallel heat flux
<code>flalfi</code> (recommend=0.21)	:for ion parallel heat flux
<code>flalfv</code> (recommend=0.5)	:for ion parallel viscosity

These parameters are all defaulted to large values (1e20 or 1e10) which gives effectively no flux limiting. The recommended values are obtained from fits to Monte Carlo and Fokker Planck calculations, but are not universal.

There are also flux-limiting coefficients for the diffusive gas fluxes. These are called `flalfgx`, `flalfgy`, and `flalfgxy` for diffusion driven by the density gradients in the x-, y-, and nonorthogonal xy- directions. In addition, fluxes can be driven by gradients in temperature which have flux-limit coefficients `flalftgx` and `flalftgy`. Finally, the neutral gas viscosity coefficients can be limited through the parameters `flalfvgx` and `flalfvgy`. All of the gas flux-limit parameters are defaulted to large numbers, but physically reasonable values are in the range of unity which

are left as an option for the user. Our experience is that the code can have difficulty with these flux limits if gradients become too steep - this is the reason for having them switched off as a default.

10.2. Extensions to flux-limit models

Two variations from the simple flux-limit model shown in the manual are described here: (1) including the effect of finite domain size such as a nearby wall or simply be an aid to improve numerical robustness of the flux-limited algorithm, and (2) dividing the limited flux into a sum of diffusive and convective terms.

Adding a finite-domain parameter

For the first extension, consider, as an example, an original diffusive particle flux $\Gamma_0 = D_0 \nabla n$ arising from collisions with other particles. Here the diffusion coefficient is expressed as

$$D_0 \sim \lambda_c^2 / \tau_c = v_t^2 \tau_c \quad (7)$$

where $\lambda_c = v_t \tau$ is a collisional mean-free path, τ_c is the collision time, and v_t is the thermal velocity. The typical flux-limited expression for the particle flux is

$$D_0 \rightarrow D_0 / [1 + |\Gamma_0 / \Gamma_l|^{\gamma_{fl}}]^{1/\gamma_{fl}} \quad (8)$$

where γ_{fl} (= `flgamg` in UEDGE) is some integer (e.g., 1 or 2), $\Gamma_l = C_l v_t n$, and C is some coefficient. The ratio

$$\Gamma_0 / \Gamma_l = (v_t^2 / \tau L_n) n / (C_l v_t n) = \frac{1}{C_l} \frac{\lambda_c}{L_n} \quad (9)$$

where $L_n \equiv |n / \nabla n|$ is the density scale length. However, if there is a wall present at some distance L_w , a heuristic generalization of the ratio Γ_0 / Γ_l is

$$\Gamma_0 / \Gamma_l = \frac{1}{C_l} \frac{\lambda_c}{L_n} (1 + L_n / L_w). \quad (10)$$

Note that there is also a potential numerical benefit to including L_w : when λ_c is large, but gradients are very small or pass through zero (L_n large), the term Γ_0 / Γ_l without L_w can rapidly change from being very large to passing through zero. This strong nonlinear behavior

can cause undesirable numerical sensitivity. The addition of a L_w restricts the variation of Γ_0/Γ_l , yielding a more robust algorithm.

UEDGE allows specifying the length L_w for an number of transport quantities independently, though the heuristic argument suggest these should be the same. Nonetheless, the independence of the different L_w allow one to investigate improved robustness of the numerical algorithm for each quantity. The UEDGE transport process and associated "wall" lengths (L_w) are as follows:

lgmax	neutral density
lgvmax	neutral parallel viscosity
lgtmax	neutral thermal conductivity
lxtemax	poloidal electron thermal conductivity
lxtimax	poloidal ion thermal conductivity

Alternate flux-limit expression

There is a another option for flux-limiting the poloidal electron and ion thermal conductivity that represents the flux as a sum of diffusion and convection. If we consider a thermal diffusivity, κ_0 , then this representation leads to the following expression for the heat flux

$$-\kappa_0 \nabla_x T \rightarrow \frac{-\kappa_0 \nabla_x T}{(1 + q_r)^2} + C_{lt} n v_t T \frac{q_r^2}{(1 + q_r)^2} \text{sign}(q_0) \quad (11)$$

where $q_r = q_0/q_l$, $q_0 = -\kappa_0 \nabla_x T$, $q_l = C_{lt} n v_t T$, and C_{lt} is the thermal flux-limit coefficient (the counterpart of C_l for neutral particle flux above). Thus, if $q_r \ll 1$, one obtains the classical heat flux, and if $q_r \gg 1$, one obtains the flux-limited heat flux q_l . The difference with the original flux-limited expression is that here the heat-flux limit appears explicitly as a convective term rather than simply a reduction of the diffusive flux, which might have some numerical advantages.

In UEDGE, the alternate thermal flux-limit expression is controlled by the input parameters `isflxle` and `isflxdi` for electrons on ions. If `isflxle,i` are 1, the original diffusive heat flux limit expressions are used, whereas if `isflxle,i` are 0, the second diffusive/convective expressions are used. The default for UEDGE is to use the alternate flux-limit expression

for electrons and the original one for ions. It should be noted that the default for ions is really `isflxldi = 2`, which means that the original flux-limit expression is used, but the individual thermal flux limited conductivities are computed for all of the individual ion and neutrals species and then summed, rather than flux-limiting the summed thermal conductivities. Also, the "wall" distance parameters discussed in the first half of this note do work for the alternate flux-limit expression as well as the original expression.

11. Models for Neutral Gas

11.1. *Inertial fluid and diffusive models for atoms*

UEDGE uses two models for the neutral gas, the most general being an inertial fluid model that solves the parallel momentum equation along the direction of the magnetic field, B , and diffusion in the two directions perpendicular to B (implemented by F. Wising). For this model, set `isupgon(1)=1`, `isngon(1)=0`, `nhsp=2`, and `ziin(2)=0`. Neutral viscosity and thermal conduction is included using both charge-exchange collisions and neutral-neutral collisions.

The simpler gas model solves a diffusion equation in the 2-D poloidal plane. For hydrogen, it is activated by setting `isngon(1)=1`, `isupgon(1)=0`, and `nhsp=1`. The diffusion coefficient is given by $D_g = T_i / [m_i * (nucx + nuiz)]$.

For both of the gas models just mentioned, one can turn off the contribution of the gradient of the neutral (ion) temperature for the velocity calculation (from the pressure gradient term in the momentum equation) by setting `cngfx(1)=0` and `cngfy(1)=0`. The default is for these terms to be active. By setting these factors for indices `igsp=2` and above, one can likewise modify the temperature gradient term for the impurity gas.

The impurity gas is modeled by the diffusion equation just mentioned and is activated by setting `isngon(2)=1` (if more than one impurity species is present, then `isngon(3)=1`, etc.). Here the diffusion coefficient is somewhat more general by including elastic collisions as $D_g = T_g / (m_g * nuix)$, where

$$\text{nuix} = \text{rcxighg} * \text{nucx_h} + \text{nuiz} + \text{nucx_imp} + \text{massfac} * (\text{kelighi} * \text{ni_h} + \text{kelighg} * \text{ng_h} + \text{keligii} * \text{ni_imp})$$

Here, `nucx_h` is the charge-exchange frequency between hydrogen neutrals and impurity ions, `rcxighg` is a scale factor (usually small) to convert this rate to the frequency of impurity neutrals C-X with hydrogen ions. Likewise, `nucx_imp` is charge-exchange of impurity neutrals with impurity ions, where the cross-section is given by `sigcxms`. To account for elastic collisions of impurity gas with hydrogen ions and gas, and finally with impurity ions, we use the last three terms. Details of this model were suggested by S. Krasheninnikov, where

$$\text{massfac} = 16 * \text{mi_h} / [3 * (\text{mi_imp} + \text{mi_h})]$$

and `kelighi` and `kelighg` are the $\langle \sigma v \rangle$ rates. Estimated values are `kelighi = kelighg = 5e-16 m**3/s` at temperatures of ~ 1 eV, but the temperature dependence is neglected. The precise values are uncertain. Values of `kelighi` and `kelighg` should be set in the users input file. Thus, if you set `rcxighg=0.`, you should set `kelighi(igsp) = kelighg(igsp) = 5.e-16`, or some more accurate value if available. This procedure prevents `D_g` from becoming very large in low temperature (~ 1 eV) regions when `rcxighg=0` and `nuiz` are very small.

11.2. Options for temperature of neutrals

One can let the neutrals have a multiple of the common ion temperature or use a specified mixture of constant value, `tgas` and the ion temperature `ti`. These options are controlled as follows (the more common two extremes are mentioned first, then the more general mixture):

If `istgcon=0`, then

the gas temperature, `tg`, is a multiple of the ion temperature, `ti`.

Specifically, then `tg=rtg2ti*ti` across the whole mesh.

If `istgcon=1`, then

`tg = tgas*ev`, where `tgas` is an input variable in eV, and `tg` has this same constant value across the mesh.

More generally, if `istgcon` is between 0 and 1 (`istgcon` is a real), then

$$tg = (1-istgcon)*rtg2ti*ti + istgcon*tgas*ev$$

If the neutral temperature, `tg`, is chosen to be different than the ion temperature, one can activate the collisional cooling of `ti` from C-X and elastic collisions by setting the factor `cftiimp` to unity; the default is `cftiimp = 0`, *i.e.*, no cooling. Similarly, the drag on the parallel velocity of impurity ions with impurity neutrals from C-X and elastic scattering is switched on by setting `cfupimp`=1.

11.3. *Inclusion of fluid molecules via the diffusive approximation*

A diffusive-neutral fluid component (*i.e.*, neglecting parallel inertia) can be used to represent the molecules which evolve from the wall to describe the thermal desorption phase of recycling; this is usually the dominant recycling channel.

To include hydrogen molecules, one should set the following input parameters:

```

ngsp = 2          # if no impurities are present
nhgsp = 2         # tells code that two hydrogen gas species are present
ishymol = 1       # switch to turn on hydrogen molecules
recycp(1)= -0.5   # neg. recycp(1) acts like -albedo for atomic gas
recycp(2)= 0.98   # recycling into molecular channel for ions + atoms
recycw(1)= -0.5   # neg. recycw(1) acts like -albedo for atomic gas
recycw(2)= 0.98   # recycling into molecular channel for ions + atoms
cdifg(2) = 0.05   # reduces mol gas diff coeff to simulate wall temp
isngon(2) = 1     # turns on the second gas species as diffusive
isupgon(2) = 0    # be sure inertial switch is off for molecules

```

Note that `recycp` applies to the divertor plates and `recycw` applies to the side walls when `matwsi,o > 0`. The second gas species then corresponds to the molecules while the first is the atomic species. Even though the molecules presently must use the diffusive approximation,

the atomic species can be either diffusive neutrals or 1-D Navier-Stokes as described just above. The diffusion coefficient for the molecular gas is $D_g = c_d \text{fig}(2) * T_g / (m_g * \nu_{ix})$, where ν_{ix} is now given by

$$\nu_{ix} = \nu_{\text{diss}} + \text{massfac} * (\text{kelighi} * n_{i_h} + \text{kelighg} * n_{g_h})$$

where ν_{diss} is the dissociation rate calculated using a polynomial fit obtained from the EIRENE neutral Monte Carlo code and **massfac** is defined in the previous section.

11.4. *Coupling to Monte Carlo neutral codes*

The hydrogenic fluid neutrals model can be turned off and replaced by a Monte Carlo neutrals model. In the simplest scheme, one uses a numerically explicit time-dependent coupling of plasma and neutrals models, with the models communicating via disk files. The UEDGE plasma mesh information is written to a disk file, **fort.30**, with the command:

```
call write30("fort.30", runid)
```

The UEDGE plasma background information, i.e., density, temperatures and flow velocity, is written to a disk file, **fort.31**, with the command:

```
call write31("fort.31", runid)
```

where **runid** is some header text to identify the run.

The hydrogenic fluid neutrals model is turned off and the Monte Carlo neutrals on via the switches:

```
nhsp=1
isupgon(1)=0
isngon(1)=0
ismcnon=1
```

which turns off the hydrogenic fluid neutrals contributions to the plasma source terms for ion density, ion parallel momentum, electron temperature and ion temperature. The user must replace these sources with corresponding sources from the Monte Carlo neutrals model at each time step before executing the plasma model with `exmain`. For the EIRENE Monte Carlo code, the procedure is as follows:

1. call `read32("fort.32")`; this reads a data file, `fort.32`, which contains normalized plasma source terms due to each 'stratum' in EIRENE; the data arrays are `sni`, `smo`, `see`, `sei` and the normalization constant(s) `wsor`.
2. convert from normalized source terms to physical source terms, e.g., `mensor_ni = -wsor*sni`
3. convert the source for total ion energy to a source for thermal ion energy only, i.e., `mensor_ti = mensor_ti - up*mensor_up + (.5*mi*up**2)*mensor_ni`
4. compute total plasma sources by summing over all 'strata', e.g., `uesor_ni(ix,iy,ifld) = sum on istra [mensor_ni(ix,iy,ifld,istra)]`

After executing the plasma model, one writes the plasma background data for the next Monte Carlo neutrals calculation with a call to subroutine `write31` as noted above.

The Monte Carlo code can be executed from within UEDGE via the parser command:

```
basisexe("eiobjx < input.dat > eir.log")
```

where `eiobjx` is the name of the executable and `input.dat` and `eir.log` are standard input and output files for the EIRENE code. Some diagnostic output from the EIRENE code is accessible within UEDGE with the parser command:

```
call read44("fort.44")
```

In particular, the atomic neutral density in the array `naf(1:nx,1:ny,1)` may be compared with the fluid model result `inng(1:nx,1:ny,1)`. A similar procedure is followed for coupling to the DEGAS2 Monte Carlo code [31].

12. Models for Hydrogen Ionization, Radiation, and Recombination

12.1. Basic rate data tables available

The calculation of the ionization, radiation, and recombination terms in the ion and gas continuity equations is taken either from an analytic model or linear interpolation of data from table look-up. For most of the options, recombination is switched on by `isrecmon=1`. The model used is controlled by the variable `istabon` and has the following options:

<code>istabon=0</code>	:Analytic model for ionization (from Braams in B2) and :charge-exchange; no recombination
<code>istabon=1</code>	:Tables from ADPAK by Hulse via Braams; not recommended :by Doug Post for the low temperature (< 50 eV) regime
<code>istabon=2</code>	:Tables from STAHL by Behringer at Garching via Braams
<code>istabon=3</code>	:Tables used in DEGAS from Janev, Post, et al., 1984
<code>istabon=4</code>	:Extended-DEGAS tables from Doug Post '93 for temperatures :down to .063 eV and densities up to $1.0\text{e}+23 \text{ /m}^3$; :linear interpolation is done for <code>rsa</code> vs $\log T_e$ and $\log n_e$; :analytic model for charge-exchange.
<code>istabon=5</code>	:Same extended-DEGAS tables as option 4, but with spline :interpolation for $\log_{10}(\text{rsa})$ vs $\log(\text{te})$ and $\log_{10}(\text{ne})$
<code>istabon=6</code>	:Same extended-DEGAS tables as option 4, but with spline :interpolation for <code>rsa</code> vs $\log(\text{te})$ and $\log_{10}(\text{ne})$
<code>istabon=7</code>	:Campbell's polynomial fit for <code>rsa</code> vs $\log_{10}(\text{te})$ and $\log_{10}(\text{ne})$
<code>istabon=8</code>	:New DEGAS tables from D. Stotler Oct '93; separate electron :radiation loss rates due to ionization and recombination; :linear interpolation as in option 4; radiative loss rates :more accurate than option 4 for low T_e and/or large n_e .
<code>istabon=9</code>	:from Stotler at PPPL (95/07/10) using $\log(T_e)$ -sigv

istabon=10	:from Stotler at PPPL (95/07/10) using log(Te)-log(sigv)
istabon=11	:from Stotler at PPPL (04/09/01) using log(Te)-log(sigv) :for optic-thin Lyman- α ;coupl. coeff. for n=4-9 hydr. lines
istabon=12	:from Stotler at PPPL (04/09/01) using log(Te)-log(sigv) :for optic-thick Lyman- α ;coupl. coeff. for n=4-9 hydr. lines
istabon=13	:from Stotler at PPPL (04/09/01) using log(Te)-log(sigv) :for optic-thick all lines;coupl. coeff. for n=4-9 hydr. lines
istabon=14	:from H.Scott (03/12/01) using log(Te)-sigv; includes te,ne,rtau :dependence. Calculates min Ly- α optical depth, rtau, if rtauxfac=1
istabon=15	:from H.Scott (03/12/01) using log(Te)-log(sigv); includes te,ne,rtau :dependence. Calculates min Ly- α optical depth, rtau, if rtauxfac=1

The presently preferred table is `istabon=10`, which is a more complete table from that reported in Ref. 32. Note that for options `istabon = 14` or `15`, you will need the hydrogen data file `ehrtau.dat`.

12.2. *Escape-factor model for Lyman- α radiation trapping*

The hydrogen Lyman- α line radiation is very often optically thick for fusion plasmas. We treat this aspect with an escape-factor model developed and calibrated by M. Adams and H. Scott from the more detailed **CRETIN** multi-group radiation transport code. This model is accessed by using the options `istabon= 14` or `15` and by setting the parameter `rtauxfac=1`.

The escape-factor model calculates the (normalized) optical depth from an (R,Z) point to the poloidal and radial boundaries, and then uses the minimum value it assess the degree of Lyman- α trapping, and the associated modification to the hydrogen atomic rates. Specifically, the normalized optical depth (or neutral line density), `rtau`, is given by

$$rtau(R, Z) = rt_scal \int_{R,Z}^{Bdry} n_g dr, \quad (12)$$

where `rt_scal=1e-16` is a normalizing factor that has be used in constructing the data tables, n_g is the hydrogen atomic density, and dr is the path to the nearby boundary. The default

for **rtau** is to take the minimum for four easily-computed values: the forward and backward values of **rtau** in the poloidal direction and likewise in the radial direction. For this purpose, the inner and outer divertor legs are decoupled. The values of **rtau** in the poloidal direction are scaled with the parameter **rtauxfac** and those in the radial direction are scaled by **rtauyfac**.

Alternately, the user may set **rtauxfac** $\leq 0.$, and calculate the values of **rtau** at the parser. For such values of **rtauxfac**, **rtau** will not be changed by running UEDGE.

12.3. *Calculation of various atomic rates available at the parser*

One can find the value of various rate parameters by calling the appropriate function from the **BASIS** parser with specific arguments (not arrays). Note that the third argument is the normalized Lyman- α optical depth, **rtau**(ix,iy); for **istabon** ≤ 13 , one should use **rtau**(ix,iy)=0. These are functions for $\langle \sigma v \rangle$ and are as follows:

rsa [te(ix,iy),ne(ix,iy),rtau(ix,iy),0]	: $\langle \sigma v \rangle_{\text{ionization}}$ [m^{**3}/s]
rcx [ti(ix,iy),ni(ix,iy),1]	: $\langle \sigma v \rangle_{\text{charge_exchange}}$ [m^{**3}/s]
rra [te(ix,iy),ne(ix,iy),rtau(ix,iy),1]	: $\langle \sigma v \rangle_{\text{recombination}}$ [m^{**3}/s]

The radiative loss rates associated with ionization and recombination can be obtained by calling the following functions:

erl1 [te(ix,iy),ne(ix,iy),rtau(ix,iy)]	: $\text{ne} * \langle \sigma v \rangle_{\text{ioniz}} E_{\text{rad}}$ [J/s]
erl2 [te(ix,iy),ne(ix,iy),rtau(ix,iy)]	: $\text{ne} * \langle \sigma v \rangle_{\text{recom}} E_{\text{rad}}$ [J/s]

The calculated collision frequencies on the 2-D mesh are stored and can be viewed in the following variables:

nuiz (ix,iy)	:= ne * rsa , ionization frequency [1/s]
nucx (ix,iy)	:= ni * rcx , charge-exchange frequency [1/s]
nurc (ix,iy)	:= ne * rra , recombination frequency [1/s]; need isrecmon =1

The total (radiation + 13.6*ev) electron energy loss per ionization on the 2-d mesh is stored and can be viewed in the following variable:

`eeli(ix,iy)` :energy loss per ionization [Joules]

13. Model for Impurity Radiation and Transport

13.1. Fixed-fraction model

The impurity radiation for the fixed-fraction model uses a look-up table based on non-equilibrium coronal results from the **MIST** code (by R. Hulse and D. Post) for a given impurity. The impurity emissivity depends on electron temperature, charge-exchange recombination on neutral hydrogen and impurity lifetime due to convection. The impurity charge-exchange rate is evaluated using the neutral hydrogen density calculated by **UEDGE**, whereas the impurity lifetime is presently specified by the user in the 2-D array `atau(ix,iy)` [sec]. The impurity density is determined as a fraction of `ne(ix,iy)` by the user-specified array `afracs(ix,iy)` (requires an initial call to subroutine `allocate`), i.e., the impurity density = `ne*afracs`. Note that the work array `afrac (=afracs)` also exists, but be sure to set `afracs(ix,iy)` to enable proper interpolation if the mesh size is changed.

The impurity radiation is removed from the electron energy equation by setting the switch `isimpon > 0`. [For `isimpon=1`, you must explicitly read a set of impurity radiation data files with the command `read setup.nitrogen`; this option is now obsolete]. For `isimpon=2`, impurity radiation data files are automatically read when the code is executed with the `exmain` command; in this case the code looks for a file in your working directory with the default name `mist.dat`.

One typically specifies a constant fraction of impurities by typing `afrac(,)=0.01` for 1%; the impurity lifetime `atau(,)` is defaulted to 1 second, so it does not affect the calculation significantly unless the user sets a lower value. We are working to improve the impurity model so that the fraction and lifetime are calculated self-consistently from transport.

13.2. Multi-species models

One can also treat the impurities by following the densities of the individual charge states. This is done by setting `isimpon=5` where the FMOMBAL package by Steve Hirshman, ORNL, is used to calculate the friction forces between species and a mass-averaged momentum equation is solved for all the species. Another option is to set `isimpon=6` where the friction forces are determined from analytic formulae and the individual impurity parallel velocities come from the force balance equation that results when impurity inertia and viscosity are ignored. The thermal force coefficients owing to parallel gradients in the electron and ion temperatures are taken from Igithanov [33] and are similar to those given by Keilhacker [34].

It is possible to solve individual parallel momentum equations for each charge state. This requires setting the parameter `nusp.imp` to the number of impurity momentum equations, and setting `isimpon=6` again. The intra-species friction is either taken from Ref. 33 (for `isofric =0`) or from the B2 expression [3] (for `isofric =1`).

The radial transport of the impurity species is determined by the diffusion coefficients, `difni`, which may be set differently for each charge state followed (see the section on anomalous transport).

The look-up tables for impurity ionization, radiation, and recombination rates come from a computer code developed by B. Braams which writes out tables for either **ADPAK** rates [35] or **STRAHL** rates [36]. The rates that are used are controlled by the character variable `mcfilename`; its use is illustrated in the multiple-isotope example shown below.

Multiple isotopes can be followed simultaneously. Here the impurity gas is modeled using an diffusion equation. The relevant parameters for a typical input file with helium and neon are as follows (this cases assumes the first two “ion” species are hydrogen ions and hydrogen neutrals using the inertial neutral model; the impurities thus start with ion species 3, but gas species 2):

```
# Impurities
      isimpon =6           #Use force balance equation
```

```

# Impurity gas
    ngsp = 3          #total number of gas species (hydrogen+ impurities)
    isngon(2:3) = 1    #turn on impurity gas equations
    recycp(2:3) = 1.0  #plate recycling coeff of helium and neon

# Helium species
    nzsp(1) = 2        #number of helium charge states used
    minu(3:4) = 4.     #helium mass in AMU
    ziin(3:4) = iota(1:2) #charge for each state used
    znuclin(3:4) = 2    #nuclear charge for helium
    isnicore(4)= 1      # =1 for fixed core density of He++
    ncore(4)=2.e18      #density of He++ at core boundary

# Neon species
    nzsp(2) = 8        #number of neon charge states used
    minu(5:12) = 20.   #neon mass in AMU
    ziin(5:12) = iota(1:8) #charge for each state used
    znuclin(5:12) = 10  #nuclear charge
    isnicore(12)= 1     # =1 for fixed core density of Ne+8
    ncore(12)=4.e17     #density of Ne+8 at core boundary

# Specify impurity data files
    nzdf = 2           #number of impurity data files to be read
    mcfilename = ["He_rates.strahl", "Ne_rates.strahl"] #data file names

```

Note that there are now two impurity data files to be read, and that the user needs to set the variable `nzdf` to 2 to specify this.

14. Specifying Anomalous Radial Transport Coefficients

The simplest description of radial transport is a set of spatially constant diffusion coefficients for density, parallel momentum, electron energy, and ion energy. All are in units of

m^{*2}/s , and the density and momentum coefficients allow 12 (expandable) locations for 12 species. The coefficients are as follows:

difni(i)	# radial (or y-direction) density diffusion
vcony(i)	# radial convective pinch velocity
difni2(i)	# perpendicular density diffusion in "2" direction in # flux surface (perp. to y and — directions)
difpr(i)	# radial ion velocity as $difpr(1/P)(dP/dy - 1.5ndTe/dy)$
travis(i)	# radial parallel momentum diffusion
difutm(i)	# radial toroidal momentum diffusion (for potential eqn)
kye	# radial electron energy diffusion
kyi	# radial ion energy diffusion

In addition, one can introduce user-specified, spatially-dependent diffusion coefficients, or let the code calculate Bohm-like diffusion coefficient that are added to the ones above. The switch `isbohmcalc` determines which of these options is active:

- `isbohmcalc = 1` → Code fills 2-D arrays `dif_use`, `tra_use`, `kye_use`, and `kyi_use` = $Te/(16eB)$ but also multiplied by `facbni`, `facbup`, `facbee`, `facbei`.
In addition, `vy_use` may be set to a user-specified array as the pinch velocity.
(default=1) (these Bohm rates are calc. on x,y-mesh-faces)
- `isbohmcalc = 0` → Code uses whatever user initially sets for arrays `dif_use`, `dif2_use`, `tra_use`, `kye_use`, `kyi_use`, and `difutm`. In addition, `vy_use` may be set to a user-specified array as the pinch velocity.
Also, one can set `vyte_use`, `vyti_use`, and `vyup_use` to carry portions of the energy and momen. fluxes; EXPERTS ONLY, for BOUT coupling
- `isbohmcalc = 2` → Similar to `isbohmcalc=1`, except harmonic average between scaled Bohm and `difni`, `travis`, `kye`, `kyi`
- `isbohmcalc = 3` → User set coefficients where the outer midplane values are determined by input `difniv`, `difniv2`, `difprv`, `travisv`, `kyev`, `kyiv`, `difutm`, `vyte`, `vyti`, `vyup`

where all of these coefficients are radial arrays, i.e., $\text{difniv}(\text{iy})$, etc.

These coefficients are all scaled by B-field ratio $(B_{\text{midp}}/B)^{\text{inbdf}}$

The net diffusion is defined by using scale factors of the density, electron energy, and ion energy separately. The net diffusion coefficients for the `isbohmcalc=1` case are thus

```

difni      → difni + facbni *dif_use(ix,iy)
difni2     → difni2 + facbni2*dif2_use(ix,iy)
travis     → travis + facbup*tra_use(ix,iy)
kye        → kye + facbee *kye_use(ix,iy)
kyi        → kyi + facbei *kyi_use(ix,iy)
vy         → vy + vy_use(ix,iy,ifld)

```

The electron and ion energy fluxes can also be carried partially by additional user-set convective velocities, `vyte_use` and `vyti_use`, with the remained carried by `kye,i_use`; this option is for `isbohmcalc=0` and should be reserved for specialists in UEDGE/BOUT coupling. Likewise, `vyup_use` can be used to convect parallel ion momentum radially.

For `isbohmcalc=1`, all of the diffusion coefficients (`dif_use`, *etc.*), are internally calculated as equal to the Bohm rate of $T_e/16eB$, BUT the pinch term, `vy_use` is left as preset by the user (`vy_use=0` is the default).

Nearly the same relation holds for `isbohmcalc=0`, except that all of the “`facb...`” factors are unity, and the user-specified values of both the diffusion and pinch coefficients are used (the `..._use` terms). Note that `difni2` and `dif2_use` typically add small contributions to the poloidal transport that is dominated by the projection of the parallel transport; we thus usually leaved `difni2 = facbni2 = dif2_use = 0`. Here `facbni`, etc. are only scalars, thus applying equally to all ion species. If one wishes to use only user-specified diffusion coefficients, be sure to set `difni`, `difni2`, `kye`, and `kyi` to zero as all but `difni2` are defaulted to unity if they are not set in the input file. Also, if the mesh size changes and the user is specifying the values of `dif_use`, etc. (`isbohmcalc=0`, `facbni=1.`), the arrays `dif_use`, etc. must be refilled after an allocate for the new mesh size is done.

15. Converting solutions from Full-Space to Half-Space & Vice Versa

It is convenient to use the ability of **BASIS** to manipulate arrays to convert a full-space solution to a half-space solution. Here we show how to do this for a single-null solution and for a lower double-null solution. Below this is a **BASIS** script that does the reverse, i.e., takes a half-space solution and symmetrizes it as the initial guess of a full-space solution. Here a full-space problem is one with two divertor plates, one at each end of the x (poloidal) domain, and a half-space problem is one where on end of the x domain is a symmetry plane. Generally, the symmetry plane is at the left boundary (`isfixlb=2`), but can also be at the right boundary (`isfixrb=2`).

15.1. Converting from full-space to half-space

Note that such conversions take place most straightforwardly if one uses fixed temperature core boundary conditions as the input power does not then need to be adjusted. If you are running with power boundary conditions (`iflcore=1`), first determine the core temperatures, set `tcoree` and `tcorei` to these and change to `iflcore=0`, and then do the conversion to a different sized space.

First, the bottom double-null case:

(Lines beginning with `#` are comments and may be omitted. Also, the letter variables are chameleon variables in **BASIS** that take on the properties of the variable to which they are set)

```
# save original solution in temporary storage:
```

```
$n=nis;$u=ups;$e=tes;$i=tis;$g=ngs;$p=phis
```

```
# set switches to do outer quadrant only:
```

```
isfixlb=2
```

for a single-null case, set

```
nxomit=nxleg(1,1)+nxcore(1,1)
```

or, for a lower double-null case, set

```
nxomit=nxleg(1,1)+nxcore(1,1)+1
```

get very crude index-interpolated solution for outer quadrant only # NOTE: do not try to converge from this state: real oldftol=ftol; ftol=1e10; exmain; ftol=oldftol # copy original solution to restart arrays:

```
nis=$n(nxomit:nxm+1,,)
```

```
ups=$u(nxomit:nxm+1,,)
```

```
tes=$e(nxomit:nxm+1,)
```

```
tis=$i(nxomit:nxm+1,)
```

```
ngs=$g(nxomit:nxm+1,,)
```

```
phis=$p(nxomit:nxm+1,)
```

For a single-null case:

Same as above, but also set the left-hand symmetry boundary conditions:

```
nis(nxomit,,) = nis(nxomit+1,,)
```

```
ups(nxomit,,) = 0.
```

```
tes(nxomit,) = tes(nxomit+1,)
```

```
tis(nxomit,) = tis(nxomit+1,)
```

```
ngs(nxomit,,) = ngs(nxomit+1,,)
```

```
phis(nxomit,) = phis(nxomit+1,)
```

```
exmain      # outer-quadrant-only solution, on original mesh; should converge easily
```

```
read doublep # script to double nxleg, nxcore, nycore, nysol
```

```
exmain      # run and hopefully converge on doubled poloidal mesh
```

15.2. Converting from a half-space to a full-space

Assumes that $\text{nxleg}(1,1)=\text{nxleg}(1,2)$ and $\text{nxcore}(1,1)=\text{nxcore}(1,2)$. Unlike the previous example, this does not use chameleon variables, but could

Set up needed “ss” work arrays:

```
real8 niss(0:2*nx+1,0:ny+1,1:nisp); real8 upss(0:2*nx+1,0:ny+1,1:nusp)
real8 tess(0:2*nx+1,0:ny+1); real8 tiss(0:2*nx+1,0:ny+1)
real8 ngss(0:2*nx+1,0:ny+1,1:ngsp); real8 phiss(0:2*nx+1,0:ny+1)
```

Fill work arrays:

```
do ix = 0, nx
  niss(ix,,1:nisp) = nis(nx+1-ix,,1:nisp)
  upss(ix,,1:nusp) = -ups(nx-ix,,1:nusp)
  tess(ix,) = tes(nx+1-ix,)
  tiss(ix,) = tis(nx+1-ix,)
  ngss(ix,,1:ngsp) = ngs(nx+1-ix,,1:ngsp)
  phiss(ix,) = phis(nx+1-ix,)
enddo
do ix = nx+1, 2*nx+1
  niss(ix,,1:nisp) = nis(ix-nx,,1:nisp)
  upss(ix,,1:nusp) = ups(ix-nx,,1:nusp)
  tess(ix,) = tes(ix-nx,)
  tiss(ix,) = tis(ix-nx,)
  ngss(ix,,1:ngsp) = ngs(ix-nx,,1:ngsp)
  phiss(ix,) = phis(ix-nx,)
enddo
```

Modify switches, allocate proper space for save variables and fill them.

```
nxomit = 0
```

```

isfixlb = 0
allocate
  nis = niss
  ups = upss
  tes = tess
  tis = tiss
  ngs = ngss
  phis = phiss

```

At this point, you may begin the run with an `exmain` command, or save the “s” variables for a restart by creating a PFB save-file, e.g.,

```
create pf_somename;write nis,ups,tes,tis,ngs,phis;close
```

ACKNOWLEDGMENTS

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under Contract DE-AC52-07NA27344.

We are pleased to acknowledge the contributions of many people to various aspects of the `UEDGE` code. J.L. Milovich was primarily responsible for constructing initial versions of the code that utilized many aspects of the finite-difference-stencil subroutines from the B2 code by B.J. Braams, while the solution algorithm differs substantially from that in B2. Other important contributors are G.R. Smith, F. Wising, P.N. Brown, A.C. Hindmarsh, R.B. Campbell, S.P. Hirshman, D.A. Knoll, S.I. Krasheninnikov, A.Yu. Pigarov, G.D. Porter, X. Bonnin, D. Coster, D. Elder, T. Epperly, T.B. Kaiser, L.L. LoDestro, D. Stotler, and A.G. Taylor. We have also benefited from numerous discussions with R.H. Cohen, A.S. Kukushkin, D.E. Post, V. Rozhansky, D.D. Ryutov, R. Schneider, and X.Q. Xu.

Appendix A. Equations used for the UEDGE code

This section closely parallels the discussion in Ref. 25, and the interested reader may find it helpful to consult that paper for some more details. The basic form of the transport equations, without cross-field drifts, corresponds to that implemented in a number of edge transport codes being based on classical fluid equations, one of the first being the B2 [3] code as later specified in the published Ref. 4. Here we give the explicit form of the equations solved in UEDGE and show the modifications UEDGE when classical cross-field drifts are included. An accompanying PDF file called `ue_geom_fig_only.pdf` shows the coordinates used in this appendix.

The dynamical fluid equations for particle continuity, parallel (to the B-field) ion velocity, and separate ion and electron temperatures are given below. The perpendicular ion velocities come from both algebraic equations involving other variables (classical drifts) and diffusion/convection coefficients that are user-specified to mimic anomalous (turbulence-driven) ion/electron fluxes (*i.e.*, the turbulence fluxes are assumed ambipolar, *i.e.*, giving equal ion and electron fluxes). The total velocity is denoted by \mathbf{u} , the classical components use the symbol v , and they differ by the (anomalous) diffusive term. Note that parallel velocity is taken to be classical, *i.e.*, $u_{\parallel} = v_{\parallel}$ with the corresponding momentum equation given below.

For the poloidal ion velocity the various components are

$$u_{ix} = \frac{B_x}{B} v_{i\parallel} + v_{x,E} + v_{ix,\nabla B}, \quad (\text{A1})$$

where the second term is the $\mathbf{E} \times \mathbf{B}/B^2$ drift and third term is the sum of the curvature and ∇B drifts scaling in tokamaks as qT_i/RB ; here q is the ion charge and R is the major radius of the tokamak. Note that we only include those drift terms giving a finite divergence of plasma fluxes since they will appear inside divergence terms in the transport equations.

For the radial ion velocity

$$u_{iy} = -\frac{D_a}{n_i} \frac{\partial n_i}{\partial y} + V_a + v_{y,E} + v_{iy,\nabla B} + v_{iy,vis}, \quad (\text{A2})$$

where D_a and V_a are the anomalous transport coefficients characterizing turbulence-driven transport (assumed ambipolar). The third and fourth terms in Eq. A2 again are the radial

components of the cross-field drifts as in Eq. A1, and the last term is an anomalous viscous drift (non-ambipolar) that gives a connection between the electrostatic potential on neighboring magnetic flux surfaces. A more detailed discussion of the cross-field drifts terms used is given in Rognlien *et al.*, Phys. Plasmas **6** (1999) 1851.

The electron velocities have just the same form as the ion velocities including terms from anomalous diffusion/convection as well as the cross-field drifts and thus will not be repeated explicitly here. Two differences from the ion velocities is in the scaling of the ∇B drift which here is $-eT_e/RB$, and the neglect of electron perpendicular viscosity owing to their much smaller gyroradii; both of these differences lead to finite current.

Turning now to the dynamical equations for the fluid variables, the ion continuity equation is

$$\frac{\partial}{\partial t}n_i + \frac{1}{V}\frac{\partial}{\partial x}\left(\frac{V}{h_x}n_i u_{ix}\right) + \frac{1}{V}\frac{\partial}{\partial y}\left(\frac{V}{h_y}n_i u_{iy}\right) = \langle\sigma_i v_{te}\rangle n_e n_g - \langle\sigma_r v_{te}\rangle n_e n_i. \quad (\text{A3})$$

The terms $\langle\sigma_i v_e\rangle$ and $\langle\sigma_r v_e\rangle$ are rate coefficients for ionization and recombination, respectively. The metric coefficients are $h_x \equiv 1/||\nabla x||$, $h_y \equiv 1/||\nabla y||$, and $V = 2\pi R h_x h_y$ is the volume element for toroidal geometry with major radius R [4]. For brevity of presentation, the metric coefficients are suppressed in the remaining equations.

The neutral atom continuity equation is (where subscript g labels hydrogen atoms)

$$\frac{\partial}{\partial t}n_g + \frac{\partial}{\partial x}(n_g v_{gx}) + \frac{\partial}{\partial y}(n_g v_{gy}) = -\langle\sigma_i v_{te}\rangle n_e n_g + \langle\sigma_r v_{te}\rangle n_e n_i. \quad (\text{A4})$$

where the source term on the RHS is the negative of the corresponding ion source.

The ion parallel momentum equation is

$$\begin{aligned} \frac{\partial}{\partial t}(m_i n_i v_{i||}) + \frac{\partial}{\partial x}\left(m_i n_i v_{i||} u_{ix} - \eta_{ix} \frac{\partial v_{i||}}{\partial x}\right) + \frac{\partial}{\partial y}\left(m_i n_i v_{i||} u_{iy} - \eta_{iy} \frac{\partial v_{i||}}{\partial y}\right) \\ = \frac{B_x}{B} \left(-\frac{\partial P_p}{\partial x}\right) - m_i n_g \nu_{cx}(v_{i||} - v_{g||}) \\ + m_g n_g \nu_i v_{g||} - m_i n_i \nu_r v_{i||} \end{aligned} \quad (\text{A5})$$

where $P_p = P_e + P_i$, $\eta_{ix} = (B_x/B)^2 \eta_{||}$ is the classical viscosity, $\eta_{iy} = m_i n \Upsilon_a$ is anomalous. The inertialess electron momentum equation has been used to eliminate the parallel electric

field and the ion-electron friction term, which are replaced by the P_e component of P_p . Also, n_g is the hydrogen neutral density, $\nu_{cx} = n_i \langle \sigma_{cx} v_{ti} \rangle$ is the hydrogen ion-neutral charge-exchange frequency, and $u_{g\parallel}$ is the (atomic) neutral hydrogen parallel velocity. All classical viscosities and thermal conductivities are flux-limited to prevent unphysically large values in regions with long mean-free paths. Expressions for the classical terms can be obtained from Ref. 1.

The corresponding parallel hydrogen atom neutral momentum equation is (with neutral variables denoted by subscript g)

$$\begin{aligned} \frac{\partial}{\partial t}(m_g n_g v_{g\parallel}) + \frac{\partial}{\partial x} \left(m_g n_g v_{g\parallel} u_{gx} - \eta_{gx} \frac{\partial v_{g\parallel}}{\partial x} \right) + \frac{\partial}{\partial y} \left(m_g n_g v_{g\parallel} u_{gy} - \eta_{gy} \frac{\partial v_{g\parallel}}{\partial y} \right) \\ = \frac{B_x}{B} \left(-\frac{\partial P_g}{\partial x} \right) + m_i n_g \nu_{cx} (v_{i\parallel} - v_{g\parallel}) \quad (\text{A6}) \\ - m_g n_g \nu_i v_{g\parallel} + m_i n_i \nu_r v_{i\parallel} \end{aligned}$$

where η_{gx} and η_{gy} are the viscosities determined by the effective collisions frequencies. In the parallel direction, the charge-exchange collisions are included as the direct momentum-exchange term proportional to $n_g \nu_{cx} (v_{i\parallel} - v_{g\parallel})$ and thus don't contribute to the parallel viscosity. However, because the poloidal magnetic field is generally much smaller than the toroidal magnetic field, the poloidal viscosity is dominated by the projection of the binormal viscosity, where one can use the charge exchange frequency, and similarly for the radial direction. Thus, ignoring the small difference between the poloidal and binormal coordinates, the viscosities in the neutral momentum equation (and the corresponding kinematic neutral thermal diffusivities below) are proportional to the kinematic diffusion coefficient, $D_g = C_v T_g / (m_g \nu_{cx})$ with T_g being the neutral gas temperature and C_v a coefficient of order unity. One then obtains $\eta_{gx} = \eta_{gy} = n_g m_g D_g$. Note that $m_i = m_g$ for the case of atomic hydrogen neutrals being considered. Viscosities are again flux-limited to prevent unphysical large values in regions with long mean-free paths.

As for the ions, the neutral velocity has two main components, one from the parallel motion just described, and the other from charge-exchange and ionization in the directions perpendicular to \mathbf{B} . Here the continuity equation is used in the inertial divergence term, which removes ionization from the neutral momentum equation. Ignoring the remaining

inertial term leads to the diffusive velocity equation:

$$\mathbf{v}_{\perp g} = \mathbf{v}_{\perp i} - \frac{\nabla_{\perp}(n_g T_g)}{m_g(n_g n_i \langle \sigma_{cx} v_{ti} \rangle + n_e n_i \langle \sigma_r v_e \rangle)}. \quad (\text{A7})$$

Thus, the total neutral velocities are

$$u_{gx} = \frac{B_x}{B} v_{i\parallel} - \frac{B_z}{B} v_{gw} \quad (\text{A8})$$

where B_z is the toroidal component of the B-field, and v_{gw} is the perpendicular component of $\mathbf{v}_{\perp g}$ in the binormal direction ($\hat{\mathbf{i}}_{\parallel} \times \hat{\mathbf{i}}_y$), For the radial ion velocity

$$u_{gy} = \hat{\mathbf{i}}_y \cdot \mathbf{v}_{\perp g} \quad (\text{A9})$$

Extension for neutral molecules

The continuity equation for neutrals can have source and sink terms; here we used an ionization term as an illustration (dissociate for molecules is similar)

$$\frac{\partial}{\partial t} n_m + \nabla \cdot (n_m \mathbf{v}_m) = -\nu_{diss} n_m. \quad (\text{A10})$$

Considering dissociation and elastic scattering source terms

$$\frac{\partial}{\partial t} n_m \mathbf{v}_m + \nabla \cdot (n_m \mathbf{v}_m \mathbf{v}_m) = -\nabla P_m / m_m - \nu_{diss} n_m \mathbf{v}_m - \nu_{ela} n_m (\mathbf{v}_m - \mathbf{v}_a) - \nu_{eli} n_m (\mathbf{v}_m - \mathbf{v}_i). \quad (\text{A11})$$

where $\nu_{ela,i} = n_{a,i} \langle \sigma_{ela,i} v_t \rangle$ are the elastic collision frequencies with atoms and ions. Expanding the left-hand side and substituting Eq. (1) for the second and third terms in Eq. (3) yields

$$\frac{\partial}{\partial t} \mathbf{v}_m + \mathbf{v}_m \nabla \cdot \mathbf{v}_m = -\nabla P_m / (m_m n_m) - \nu_{ela} (\mathbf{v}_m - \mathbf{v}_a) - \nu_{eli} (\mathbf{v}_m - \mathbf{v}_i). \quad (\text{A12})$$

Note that the dissociation term has dropped out of this final molecular velocity equation. Neglecting the inertia terms (LHS) for subsonic flow, yields

$$\mathbf{v}_m = -\nabla P_m / (m_m n_m \nu_{elai}) + (\nu_{ela} / \nu_{elai}) \mathbf{v}_a + (\nu_{eli} / \nu_{elai}) \mathbf{v}_i. \quad (\text{A13})$$

where $\nu_{elai} = \nu_{ela} + \nu_{eli}$.

The electron energy equation is

$$\begin{aligned} \frac{\partial}{\partial t} \left(\frac{3}{2} n_e T_e \right) + \frac{\partial}{\partial x} \left[C_{ex} n_e u_{ex} T_e - \kappa_{ex} \frac{\partial T_e}{\partial x} - 0.71 n_e T_e \frac{B_x}{B} \frac{J_{\parallel}}{en_e} \right] + \frac{\partial}{\partial y} \left(C_{ey} n_e u_{ey} T_e - \kappa_{ey} \frac{\partial T_e}{\partial y} \right) \\ = \left[u_{ix} \frac{\partial P_e}{\partial x} - u_{iy} \frac{\partial P_i}{\partial y} - u_{iw} \frac{B_x}{B} \frac{\partial P_p}{\partial x} \right] + \mathbf{E} \cdot \mathbf{J} - K_q (T_e - T_i) + S_{Ee}. \end{aligned} \quad (\text{A14})$$

Here the poloidal heat conductivity is classical, $\kappa_{ex} = (B_x/B)^2 \kappa_{\parallel}$, radial is anomalous, $\kappa_{ey} = n \chi_e$, and K_q is the collisional energy exchange coefficient. The velocity u_{iw} is that in the direction binormal to \mathbf{B} and the radial direction ($\hat{\mathbf{i}}_{\parallel} \times \hat{\mathbf{i}}_y$), which is only used when cross-field drifts are included. Typical values for convection coefficients $C_{ex,ey}$ are 5/2 or 3/2.

The ion energy equation below has an implied sum over the ion and neutral species with $T_g = T_i$;

$$\begin{aligned} \frac{\partial}{\partial t} \left(\frac{3}{2} n_i T_i \right) + \frac{\partial}{\partial x} \left[C_{ix} n_i u_{ix} T_i - \kappa_{jx} \frac{\partial T_i}{\partial x} \right] + \frac{\partial}{\partial y} \left(C_{iy} n_i u_{iy} T_i - \kappa_{jy} \frac{\partial T_i}{\partial y} \right) = \left[\mathbf{u}_i \cdot \nabla P_i \right] \\ + \eta_{ix} \left(\frac{\partial v_{ij\parallel}}{\partial x} \right)^2 + \eta_{iy} \left(\frac{\partial v_{i\parallel}}{\partial y} \right)^2 + K_{qj} (T_e - T_i) + \frac{1}{2} m_i v_{i\parallel}^2 n_i \nu_{iz} + S_{Ej}. \end{aligned} \quad (\text{A15})$$

As for the electrons, the poloidal thermal conduction (and viscosity) coefficients are classical and the radial are anomalous; again $C_{ix,iy}$ are typical either 5/2 or 3/2.

The equation for the potential is obtained by subtracting the ion and electron continuity equations, and assuming quasineutrality, $n_i = n_e$:

$$\nabla \cdot \mathbf{J}(\phi) = \frac{\partial}{\partial x} (J_x) + \frac{\partial}{\partial y} (J_y) = 0 \quad (\text{A16})$$

Here by \mathbf{J} we mean the currents excluding the magnetization current since the divergence of the latter is automatically zero owing to it being the curl of a vector. The remaining current components are

$$\mathbf{J} = \left[ne(\mathbf{v}_{i,\nabla B} - \mathbf{v}_{e,\nabla B}) \cdot \hat{\mathbf{i}}_x + J_{\parallel} \frac{B_x}{B} \right] \hat{\mathbf{i}}_x + ne(v_{i,y1} - v_{e,y1}) \hat{\mathbf{i}}_y. \quad (\text{A17})$$

Note that the terms arising from the $\mathbf{v}_{\nabla B}$ -drift do not explicitly depend on ϕ , so they act as source terms in Eq. (A16). The expression for the parallel current, J_{\parallel} , comes from the parallel momentum equation for electrons with $m_e \rightarrow 0$, yielding,

$$J_{\parallel} = \frac{en}{0.51m_e \nu_e} \frac{B_x}{B} \left(\frac{1}{n} \frac{\partial P_e}{\partial x} - e \frac{\partial \phi}{\partial x} + 0.71 \frac{\partial T_e}{\partial x} \right). \quad (\text{A18})$$

Here ν_e is the electron collision frequency, and the numerical coefficients are described in Ref. 1. Note that the expression for the radial current is different from that in Ref. 4.

A basic model for the impurity parallel velocity is obtained from the ion parallel momentum equation neglecting inertia and viscosity terms

$$\begin{aligned} \frac{\partial}{\partial t}(m_z n_z v_{z\parallel}) + \frac{\partial}{\partial x} \left(m_z n_z v_{i\parallel} u_{zx} - \eta_{zx} \frac{\partial v_{z\parallel}}{\partial x} \right) + \frac{\partial}{\partial y} \left(m_z n_z v_{z\parallel} u_{zy} - \eta_{zy} \frac{\partial v_{z\parallel}}{\partial y} \right) \\ = \left(\frac{B_x}{B} \right) \left(-\frac{\partial P_z}{\partial x} + \alpha_z n_z \frac{\partial T_e}{\partial x} + \beta_z n_z \frac{\partial T_i}{\partial x} - Z e n_z \frac{\partial \Phi}{\partial x} \right) \\ - (u_{\parallel z} - u_{\parallel h}) m_h n_h \nu_{hz}. \end{aligned} \quad (\text{A19})$$

Note that the B_x/B converts the poloidal gradients into parallel gradients, i.e., $(B_x/B)\partial/\partial x = \partial/\partial s$, where s is the parallel distance along a field line. The thermal force coefficients come from M. Keilhacker et al. [Nucl. Fusion **31** (1991) 535] and Yu. Igithanov [Contr. Plasma Phys. **28** (1988) 477]:

$$\alpha_z = 2.2Z^2(1 + 0.52Z_{eff})[(1 + 2.65Z_0)(1 + 0.285Z_{eff})Z_{eff}], \quad (\text{A20})$$

$$\begin{aligned} \beta_z = 1.56(1 + \sqrt{2}Z_0)(1 + 0.52Z_0)(1 + \sqrt{2}Z_0)(1 + 0.52Z_0)Z^2(1 + \sqrt{2}Z_0)(1 + 0.52Z_0) \\ [(1 + 2.65Z_0)(1 + 0.285Z_0)\{Z_0 + ((m_h + m_z)/2m_z)^{1/2}\}] \\ + 0.6(Z^2 n_{tz}/n_h Z_0 - 1), \end{aligned} \quad (\text{A21})$$

where $Z_0 = n_e Z_{eff}/n_h - 1$, Z is the charge of the impurity ion, the h subscript refers to the hydrogen ion, and n_{tz} is the total impurity ion density summed over all charge states (with a common ion temperature).

A common option used in UEDGE is to omit the impurity inertial effects or left-hand side of Eq.(A19) since parallel velocities are typically much less than the sonic speed except near the divertor plates. This is called the force-balance model, and is activated in UEDGE by setting `isimpon = 6` and the number of parallel hydrogen momentum equations `nusp = 1` or `2` depending on whether the hydrogen neutrals are diffusive only (1) or include inertia (2).

A more complete option is to include the individual impurity ion inertial terms in Eq. (A19); this is achieved in UEDGE by setting the input variable `nusp = total num-`

ber of full ion momentum equations (generally one each for hydrogen ions and neutrals, plus the number of impurity charge states, excluding the neutral impurity).

Appendix B. UEDGE topology conventions




Figure 1: Single-null geometry of the 2D edge region simulated showing the coordinates used and the magnetic-field components; here I_p is the plasma current that generates the poloidal field B_x . The vector \mathbf{w} is in the so-called bi-normal direction $\hat{\mathbf{i}}_{||} \times \hat{\mathbf{i}}_y$ with $\hat{\mathbf{i}}_{||}$ being along the total \mathbf{B} -field.

Appendix C. A Primer on Input Files and Obtaining Initial UEDGE Solutions

A Primer on Input Files and Obtaining Initial UEDGE Solutions

Tom Rognlien
Updated 03/19/15

This document is for the user who wants to create a new UEDGE case from a different MHD equilibrium-based mesh or for a user who wants to learn a general procedure to create new cases without relying entirely on the previous cases. Nevertheless, for either situation, the user should find it useful to use the existing test case closest to the desired new case as a rough guide, although that will not be assumed here (and thus the length of this document). For users using the Python version of UEDGE from the FACETS project, there is a set of standard input files for various test cases located in the FACETS SVN directory `uedge/test/Forthon_cases` (see the file `README-FORTHON_cases` in this directory for a description of each case).

If instead, you are using a Basis version of UEDGE, the instructions given below are the same except the following: Firstly, in the Basis version, one does not need to use the prefix on variable names (e.g., `bbb.mhdgeo` can be truncated to `mhdgeo`), though you may use the prefix and Basis will understand. Secondly, all Python arrays begin from 0, whereas in Basis, only arrays explicitly dimensioned as starting at 0 will be treated this manner. Also, Python uses `[]` brackets, whereas Basis uses `()` brackets for array indexes. In the text below, any Basis-type array notation is shown at the end of the comment for each variable. Also note that in calling a subroutine with no arguments from Basis, the `()` can be omitted, e.g., `bbb.exmain()` is equivalent to `exmain`.

A more complete description of the options and features of the UEDGE code is given in the manual available as http://www.mfescience.org/mfedocs/uedge_man_V4.41.pdf.

STEP 1 - MHD equilibrium and geometry type:

Obtain the ASCII-formatted "a" and "n or g" EFIT format files for the desired MHD equilibrium; UEDGE uses information in both of these files. Here we treat the most common (and default) case of a single-null equilibrium. The following input parameters are then added to the input file:

```
bbb.mhdgeo=1          # Specifies tokamak toroidal geometry
bbb.isfixlb=0         # 0 for standard inner divertor plate BCs
bbb.isfixrb=0         # 0 for standard outer divertor plate BCs
os.system('cp -f aXXX123 aeqdsk') # aeqdsk file name; for Basis omit os.system and ()
os.system('cp -f gXXX123 neqdsk') # neqdsk file name; for Basis omit os.system and ()
```

Double null geometries can be treated as well. Here the simplest case is to consider up/down symmetry about the magnetic axis, and the only change is to add the input variable

```
bbb.geometry = "dnbot"    # Default is "snull"
```

UEDGE can also treat 2D slab geometry (`bbb.mhdgeo=0`) and cylindrical (`r,z`) geometry (`bbb.mhdgeo=-1`) with internal mesh algorithms; see the general user manual.

STEP 2 - Mesh setup:

There are many input variables that can be used to optimize the mesh cell spacing as described in the user manual; here we give a basic set for an orthogonal mesh where two of the cell faces are

aligned with the poloidal flux surfaces (only using a linear curve) and the other two are orthogonal to the first two.

The initial mesh should probably be of modest size because it is then faster and more trouble-free to obtain an initial steady-state UEDGE solution, and finer-mesh solutions can subsequently be obtained by using UEDGE's ability to interpolate a solution of one mesh to another - using essential arbitrary size ratios - as an initial guess on the finer mesh. But it will be best to consider increasing the mesh size incrementally for efficiency. A second consideration of the initial mesh size is if one has a previous solution for an even remotely related problem saved as an HDF5 (or PDB) file. Targeting this pre-existing-case mesh size can help lead to a new solution quickly. However, even this pre-existing case should preferably be of moderate size (one can use the UEDGE interpolation to decrease mesh size on the pre-existing problem), as the new case will need to begin with the mesh size of the pre-existing case. More detail is given below for determining the initial guess (starting values) for the solution.

The user controls the radial size of the simulated region by specifying the minimum and maximum values of the poloidal flux (normalized to unity at the magnetic separatrix). There are two minimum flux values used, one each for the core and private-flux regions. The poloidal boundaries are typically (but not always) given by the position of the divertor plates in the MHD equilibrium "g" files, although these can be overridden by user input. To determine the radial size, set

```
flx.psi0min1 = 0.95      #core minimum
flx.psi0min2 = 0.98      #private flux minimum
flx.psi0max = 1.05       #maximum flux at wall
```

If you want (or need) to specify the location that the separatrix intersects the divertor plates, you only need to give the intersection of the major radius (grd.rstrike) with the separatrix on the inner and outer divertor plates (thus two values); thus, set

```
grd.isspnew = 1          # Turns on user-specified strike point
grd.rstrike = [1.77, 2.5] # R_major [inner, outer] at sep intersec; for Basis (1.77,2.5)
```

To determine the mesh size, the domain is divided into four poloidal sections moving clockwise poloidally from the inner divertor plate: inner divertor leg, inner core, outer core, and outer divertor leg. Radially, there are two segments moving outward from the core boundary: core and scrape-off layer (SOL). The number of cells in each of these regions is set as follows:

```
com.nxleg[0,0] = 8      # Poloidal cells in inner leg; for Basis (1,1)
com.nxleg[0,1] = 8      # Poloidal cells in outer leg; for Basis (1,2)
com.nxcore[0,0] = 8     # Poloidal cells in inner core; for Basis (1,1)
com.nxcore[0,1] = 8     # Poloidal cells in outer core; for Basis (1,2)
com.nycore[0] = 10      # Radial cells in core/PF regions; for Basis (1)
com.nysol[0] = 16       # Radial cells in SOL region; for Basis (1)
```

The user can control the radial spacing of the cells by setting

```
flx.alfcy = 2.          # if <<1, most uniform; > 1 increase sep. concentration
```

Modest control of near-X-point poloidal cells by setting

```
grd.slpxt = 1.1         # Typical 1.0->1.2; larger giving concentr. near X-point
```

STEP 3 - Specify variables to be evolved:

In UEDGE, one can turn evolution of variables on and off through various switches. The normal default is for one hydrogenic ion species with density and parallel velocity equal to an assumed electron species (ambipolar transport), electron and ion temperatures (separate), and a diffusive neutral species. Typically, the diffusive neutral transport in the poloidal direction is augmented by also solving for a full momentum equation in that direction, which requires the settings listed below:

bbb.isnion[0] = 1	#=1 default) for evolving ion density; for Basis (1)
bbb.isupon[0] = 1	#=1 default) for evolving ion parallel momentum; for Basis (1)
bbb.isteon = 1	#=1 default) for evolving electron temperature
bbb.istion = 1	#=1 default) for evolving ion temperature
bbb.isngon[0] = 0	#=1 default) for evolving diffusive neutrals; for Basis (1)
bbb.nhsp = 2	#=1 default) is number of hydrogenic species (i+g)
bbb.ziin[1] = 0	#=1 default) is ion charge for second H species; for Basis (2)
bbb.isnion[1] = 1	#=0 default) for evolving inertial neutrals; for Basis (2)
bbb.isupgon[0] = 1	#=0 default) for evolving inertial neutrals; for Basis (1)

If any variable is turned off (=0), it is frozen at its present value.

STEP 4 - Set boundary conditions:

There are a large number of possible boundary conditions combinations as detailed in the UEDGE manual. Here we describe a standard basic set.

For the core boundary:

bbb.isnicore[0] = 1	#=1 uses ncore for density BC; for Basis (1)
bbb.ncore[0] = 6.0e19	#value of core density if isnicore=1; for Basis (1)
bbb.isupcore[0] = 0	#=0 default) sets u_par = 0; for Basis (1)
bbb.iflcore = 1	#specify core power
bbb.pcoree = 5.0e7	#electron power across core
bbb.pcorei = 5.0e7	#ion power across core
bbb.isngcore[0]=2	# use ionization scale length for gas; for Basis (1)
bbb.isupcore[1] = 0	#=0 default) sets gas u_par = 0; for Basis (2)

For the divertor plates:

isextrnp = 0	#=0 default) sets dni/dx=0;if=1, extrapolate
isupss = 0	#=0 default) sets u_par=cs;if=1, extrapolate
ibctep1 = 1	#=1 default) elec eng flux=bcee*Te*part_flux
ibctipl = 1	#=1 default) ion eng flux=bcee*Te*part_flux
bbb.recyep = 1.0	#ion recycling coeff. giving gas flux in
bbb.recyem = 0.1	#neut. up = -recyem*(ion up)

For private-flux walls:

bbb.isnwcon1 = 3	#PF wall density scale length bbb.lyni(1)
bbb.lyni[0] = 0.05	#PF wall density scale length; for Basis (1)
bbb.isupwi = 1	#=1 default) zero ni & ng par momem-dens flux
bbb.istepfc = 3	#priv. wall has Te scale length bbb.lyte(1)

```
bbb.istipfc = 3          #priv. wall has Ti scale length bbb.lyti(1)
bbb.recycw = 0          #ion recycling coeff. giving gas flux in
```

For outer walls:

```
bbb.isnwcono = 3        #outer wall has fixed density scale length
bbb.lyni[1] = 0.05      #outer wall density scale length; for Basis (2)
bbb.isupwo = 1          #(=1 default) zero ni & ng par momem-dens flux
bbb.istewc=3            #outer wall has Te scale length bbb.lyte(2)
bbb.istiwc=3            #outer wall has Te scale length bbb.lyti(2)
bbb.recycw = 0          #ion recycling coeff. giving gas flux in
```

STEP 5 – Setting radial transport coefficients:

The simplest model for radial transport is that they are spatially constant, and the user may then set the following in the input file (for Basis, omit []).

```
bbb.difni[]             # Ion density diff. coeff. for each species (all m**2/s)
bbb.travis[]            # Ion parallel velocity (radial) coefficient
bbb.kye                 # Electron energy coefficient (Chi_e)
bbb.kyi                 # Ion energy coefficient (Chi_i)
bbb.vcony[]             # Ion density radial convection velocity (m/s)
```

One can also add to these coefficients a corresponding set of spatially varying coefficients that can be specified after a call to allocate to properly dimension the arrays:

```
bbb.dif_use[ix,iy,ifld] # Ion density coeff.;3 indices are poloidal, radial, species
bbb.dif_travis[ix,iy,ifld] # Ion parallel velocity (radial) coeff.
bbb.kye_use[ix,iy]       # Electron energy coefficient
bbb.kyi_use[ix,iy]       # ion energy coefficient
bbb.vy_use[ix,iy,ifld]  # Ion density radial convection velocity (m/s)
```

The values of these arrays can be use to better-fit profiles to experimental values. An interpretive transport mode of UEDGE (using scripts) can help specify the coefficient profiles needed to fit the plasma profiles in the core region assuming that plasma variables are approximately constant on magnetic flux surfaces.

Yet another option that is used is to specify the radial variation of the coefficients at the outer midplane, and then UEDGE provides the poloidal variation as some power of $1/B_{\text{toroidal}}$ to easily simulate “ballooning” transport, i.e., a peaking at the outer midplane. To use this option, set

```
bbb.isbohmcalc = 3      # Flag to use poloidal variation of (1/B)**inbdif
bbb.difniv[iy,ifld]     # Ion density coeff.; indices are radial and species
bbb.travisv[iy,ifld]    # Ion parallel velocity (radial) coefficient
bbb.kyev[iy]            # Electron energy coefficient (Chi_e)
bbb.kyiv[iy]            # Ion energy coefficient (Chi_i)
bbb.vconyv[iy,ifld]     # Ion density radial convection velocity (m/s)
```

Note that for this option, the constant diffusion coefficients (e.g., bbb.difni,...) will be added, so set them to zero if a constant value is not also wanted.

STEP 6 - Choosing initial conditions:

To begin a new case, the user can either start from a set of generic profiles by setting `bbb.restart = 0`. However, as mentioned above, it is usually preferable to restart from an existing HDF5 or PDB solutions from another case, BUT WITH THE SAME NUMBER OF RADIAL AND POLOIDAL MESH CELLS. Even if the previous case is not from the same tokamak, this approach is usually easier. It is even better if the boundary conditions are the same or close (i.e., same core density or same core power), but this is not necessary. In this case, the input file will contain the lines

```
bbb.restart = 1
bbb.allocate()
uefacets.restore('yourfile.h5') # or 'yourfile.pdb' for a PDB "save" file; for Basis use the
line restore yourfile.pdb
```

Note: if you have a previous case that is on a different size mesh, you can use the UEDGE interpolator to produce an approximate solution on the desired mesh size by first running the prior case interactively as follows:

a) Start up by having UEDGE read the prior-case input file, but don't execute `bbb.exmain` yet (doesn't really hurt to do so, just takes time)

b) Then set the following:

```
bbb.issfon = 0      # Don't calculate initial Jacobian
bbb.ftol = 1.e50    # Allowed tolerance for an "accepted" solution
bbb.exmain()
```

c) Now change the mesh values (`bbb.nxleg`, `bbb.nxcore`, `bbb.nycore`, `bbb.nysol`) to the desired values and type (UEDGE automatically interpolates between meshes)

```
bbb.exmain()
```

d) Finally save the "accepted" solution on the new mesh into a HDF5 or PDB file for use as initial conditions in the new case

Lastly note that it is possible to set the initial UEDGE profiles manually after a call to the `allocate` routine that allocates the `nis`, `ups`, etc. arrays. At this point, one can fill these arrays interactively (i.e., from PYTHON or BASIS command line) with any profiles you want, although this does require knowledge of the various geometry regions (core, SOL, private flux) to obtain sensible profiles.

STEP 7 - Obtaining an initial solution:

After creating the desired input file as described above, it is best to develop a new steady-state solution by running UEDGE interactively. The approach taken is a conservative one, outlining a series of substeps taken if the new case is substantially different than a previous case used as a template. If the new case is similar to a previous case, the procedure can often be shortened. As the user becomes more experienced, or similar. As mentioned above, it is prudent to begin new cases with modest-sized meshes, and then interpolate to finer meshes. Granted, this procedure can be time-consuming, but it generally is the most efficient and effective strategy.

7.1) Startup UEDGE and read the input file

7.2) Set

```

bbb.dtrear = 1e-9    # UEDGE time-step (s)
bbb.itermx = 15      # Maximum number on nonlinear iterations
bbb.exmain()

```

Note: if the new case is similar to a previous case, the user may be more aggressive in choosing the initial time-step, say `bbb.dtrear = 1e-5`, `1.e-3`, or in some cases `bbb.dtrear=1e20`, in which case the steady-state solution is obtained in this step alone (the skip directly to 7.5).

7.3) If convergence is not obtain (printed item .ne. 1) then skip and go to 7.4. If convergence is obtained at this single small time-step, then run UEDGE to steady-state using

```

bbb.rundt() # for older Basis versions, use the scripts rdinitdt and rdcontdt

```

This subroutine advances the time-step progressively over a number of time-steps to reach a steady-state solution. Completing this, save the new solution - you are done with the initial case; go to 7.5.

7.4) If the initial try in 7.2 does not converge, then

```

bbb.isbcwdt = 1      # Turns on the time-step for boundary equations
bbb.dtrear = 1.e-10
bbb.exmain()

```

If UEDGE convergences, type

```

bbb.t_stop = 1.e-7
bbb.rundt() # for older Basis versions, use the scripts rdinitdt and rdcontdt

```

If this finishes, you will be at a total accumulated time of `1e-7` secs. Then type

```

bbb.isbcwdt = 0      # Return boundary condition equations algebraic
bbb.dtrear = 1.e-9
bbb.t_stop = 10.
bbb.rundt() # for older Basis versions, use the scripts rdinitdt and rdcontdt

```

As under 7.3, this will advance the solution over many (expanding) time-steps to 10 secs (accumulated), generally long enough to have a essentially a steady-state solution. Completing this, save the new solution - you are done with the initial case; go to 7.5.

If UEDGE does not converge after the initial try of 7.4, or subsequently in this section, you can try turning off (and then back on) subsets of the equations with `bbb.isnion`, `bbb.isupon`, `bbb.isteon`, and `bbb.istion`, and repeat step 7.4. However, if you reach this stage, you probably should consult with an experienced UEDGE user.

7.5) If you started on a modest-sized mesh as recommended, you probably want to increase the mesh. At first, you can probably succeed in doubling the mesh in one direction, and returning to step 7.2 - this step should go more smoothly now. Then double the mesh in the second direction and repeat. As the mesh gets finer [say, $(nx,ny) \sim (50,30)$], especially if the case has many impurity species, you may find that doubling the mesh is too aggressive, and you will need to increment more slowly.

Once you have the desired mesh-size, the new test case can be used as to give the initial conditions to efficiently survey the sensitivity of the solution to power input, gas input, core density, impurity fraction, and other parameters of interest.

References

- [1] S.I. Braginskii, Transport processes in a plasma *Reviews of Plasma Physics*, Vol. I, Ed. M.A. Leontovich (Consultants Bureau, New York, 1965), p. 205.
- [2] T.D. Rognlien and D.D. Ryutov, “Psuedoclassical Transport Equations for Magnetized Edge-Plasmas in the Slab Approximation,” *Plasma Phys. Reports* **25**, 943 (1999).
- [3] B.J. Braams, “A Multi-Fluid Code for Simulation of the Edge Plasma in Tokamaks,” NET Rept. No. 68, Jan., 1987; “Computational studies in tokamak equilibrium and transport” (thesis, Univ. Utrecht, the Netherlands, 1986).
- [4] B.J. Braams, “Radiative Divertor Modeling for ITER and TPX,” *Contrib. Plasma Phys.* **36**, 276 (1996).
- [5] P.N. Brown and A.C. Hindmarsh, “Matrix-Free Methods for Stiff Systems of ODEs,” *SIAM J. Num. Anal.* **23**, 610 (1986).
- [6] Y. Saad, *Iterative Methods for Sparse Linear Systems* (PWS Pub. Co., Boston, MA, 1996).
- [7] T.D. Rognlien, J.L. Milovich, M. E. Rensink, and G.D. Porter, “A Fully Implicit, Time Dependent 2-D Fluid Code for Modeling Tokamak Edge Plasmas,” *J. Nucl. Mater.* **196-198**, 347 (1992).
- [8] T.D. Rognlien, J.L. Milovich, M.E. Rensink, and T.B. Kaiser, “Simulation of Tokamak Divertor Plasmas Including Cross-Field Drifts,” *Contr. Plasma Phys.* **32**, 485 (1992).
- [9] G.D. Porter, M. Fenstremacher, R. Groebner, *et al.*, “Benchmarking UEDGE with DIII-D Data, *Contr. Plasma Phys.* **34**, 454 (1994).
- [10] T.D. Rognlien, P.N. Brown, R.B. Campbell, *et al.*, “2-D Fluid Transport Simulations of Gaseous/Radiative Divertor,” *Contr. Plasma Phys.* **34**, 362 (1994).
- [11] G.R. Smith, P.N. Brown, R.B. Campbell, *et al.*, “Techniques and Results of Tokamak-Edge Simulation,” *J. Nucl. Mat.* **220-222**, 1024 (1995).

- [12] M.E. Fenstermacher, *et al.*, “UEDGE and DEGAS Modeling of the DIII-D Scrape-Off Layer Plasma,” J. Nucl. Mat. **220-222**, 330 (1995).
- [13] G.D. Porter, *et al.*, “Simulation of Experimentally Achieved DIII-D Detached Plasmas Using the UEDGE Code,” Phys. Plasmas **3**, 1967 (1996).
- [14] F. Wising, D.A. Knoll, S. Krasheninnikov, and T.D. Rognlien, “Simulation of Detachment in ITER-Geometry Using the UEDGE Code and a Fluid Neutral Model,” Contr. Plasma Phys. **36**, 309 (1996).
- [15] F. Wising, D.A. Knoll, S. Krasheninnikov, T.D. Rognlien, and D.J. Sigmar, “Simulation of the Alcator C-Mod Divertor with an Improved Neutral Model,” Contr. Plasma Phys. **36**, 136 (1996).
- [16] T.D. Rognlien, B.J. Braams, and D.A. Knoll, “Progress in Integrated 2-D Models for Analysis of Scrape-Off Layer Transport Physics,” Contr. Plasma Phys. **36**, 105 (1996).
- [17] T.D. Rognlien, J.A. Crotinger, G.D. Porter, *et al.*, “Simulation of the Scrape-Off Layer Plasma During a Disruption,” J. Nucl. Mat. **241-243**, 590 (1997).
- [18] F. Wising, S.I. Krasheninnikov, D.J. Sigmar, D.A. Knoll, T.D. Rognlien, B. LaBombard, B. Lipschultz, and G. McCracken, “Simulation of Plasma Flux Detachment in Alcator C-Mod and ITER,” J. Nucl. Mat. **241-243**, 273 (1997).
- [19] T.D. Rognlien and D.D. Ryutov, “Analysis of Classical Transport Equations for the Tokamak Edge Plasma,” Contr. Plasma Phys. **38**, 152 (1998).
- [20] T.D. Rognlien, D.D. Ryutov, and N. Mattor, “Calculation of 2-D Profiles for the Plasma and Electric Field near a Tokamak Separatrix,” Czech. J. Phys. **48/S2**, 201 (1998).
- [21] M.E. Rensink, L.L. LoDestro, *et al.*, “A Comparison of Neutral Gas Models for Divertor Plasmas,” Contr. Plasma Phys. **38**, 325 (1998).
- [22] T.D. Rognlien, G.D. Porter, and D.D. Ryutov, “Influence of ExB and Grad-B Terms in 2-D Edge/SOL Transport Simulations,” J. Nucl. Mater. **266-269**, 654 (1999).

- [23] M.E. Rensink and T.D. Rognlien, “Edge Plasma Modeling of Limiter Surfaces in a Tokamak Divertor Configuration,” J. Nucl. Mater. **266-269**, 1180 (1999).
- [24] S.I. Krasheninnikov, M.E. Rensink, T.D. Rognlien, *et al.*, “Stability of the Detachment Front in a Tokamak Divertor,” J. Nucl. Mater. **266-269**, 251 (1999).
- [25] T.D. Rognlien, D.D. Ryutov, N. Mattor, and G.D. Porter, “Two-Dimensional Electric Fields and Drifts near the Magnetic Separatrix in Divertor Tokamaks,” Phys. Plasmas **6**, 1851 (1999).
- [26] M. Petravic, “Orthogonal Grid Construction for Modeling of Transport in Tokamaks,” J. Comp. Phys. **73**, 125 (1987).
- [27] G.D. Byrne, “Pragmatic experiments with Krylov methods in the stiff ODE setting,” in *Computational Ordinary Differential Equations*, edited by J.R. Cash and I. Gladwell, (Oxford Univ. Press, Oxford, U.K., 1992) p. 323.
- [28] P.N. Brown and Y. Saad, “Hybrid Krylov methods for nonlinear systems of equation,” SIAM J. Sci. Stat. Comput. **11**, 450 (1990).
- [29] A.C. Hindmarsh and A.G. Taylor, “PVMODE and KINSOL: Parallel software for differential and nonlinear systems,” Tech. Rpt. UCRL-ID-129739, Lawrence Livermore National Lab. (1998).
- [30] T.D. Rognlien, X.Q. Xu, and A.C. Hindmarsh, “Application of parallel implicit methods to edge-plasma numerical simulations,” J. Comp. Phys. **175**, 249 (2002).
- [31] D.P. Stotler, *et al.*, “Coupling of Parallelized DEGAS 2 and UEDGE Codes,” Contr. Plasma Phys. **40** 221 (2000).
- [32] D.P. Stotler, D.E. Post, and D. Reiter, Bull. Am. Phys. Soc. **38** (1993) 1919; D.P. Stotler, C.S. Pitcher, C.J. Boswell, *et al.*, J. Nucl. Mater. **290-293** 967 (2001). See also the web site <http://w3.pppl.gov/degas2/>.
- [33] Yu.L. Igithanov, Contr. Plasma Phys. **28**, 477 (1988).

- [34] M. Keilhacker, R. Simonini, A. Taroni, and M.L. Watkins, Nucl. Fusion **31**, 535 (1991).
- [35] R.A. Hulse, Nucl. Tech./Fusion **3** (1983) 259.
- [36] K. Behringer, “Description of the Impurity Transport Code STRAHL,” JET Report R(87)08 (1987).